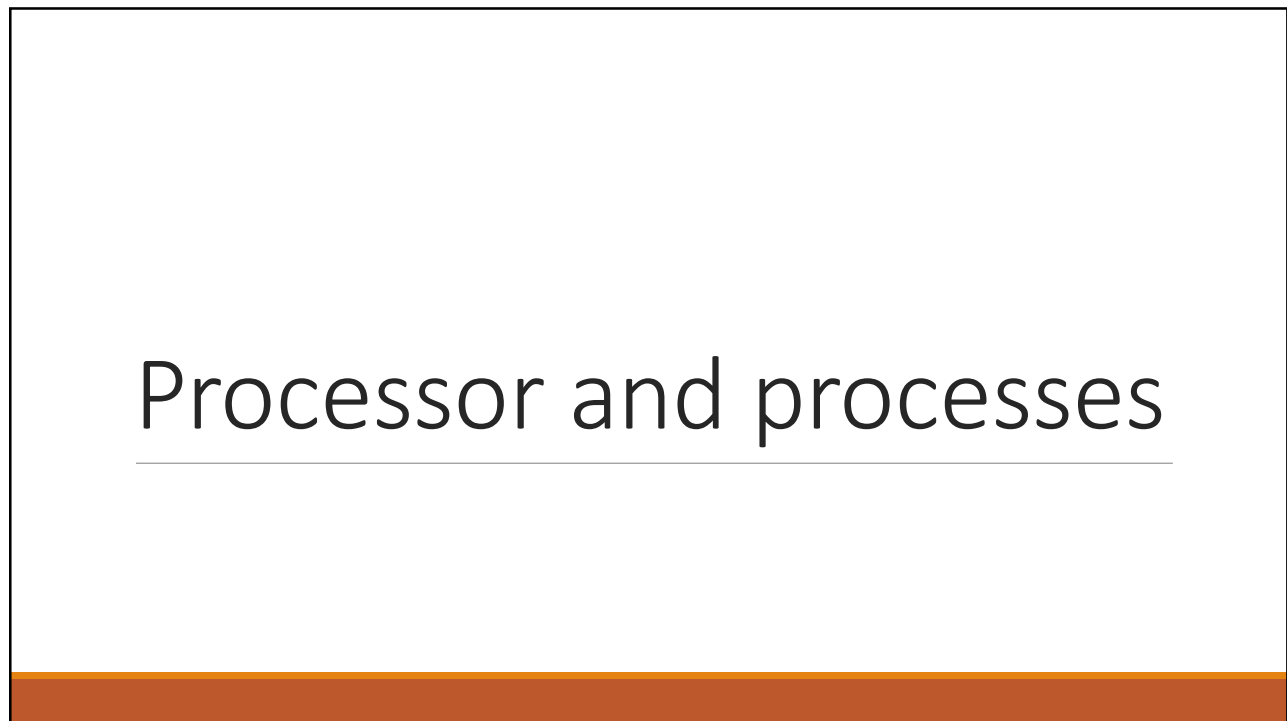
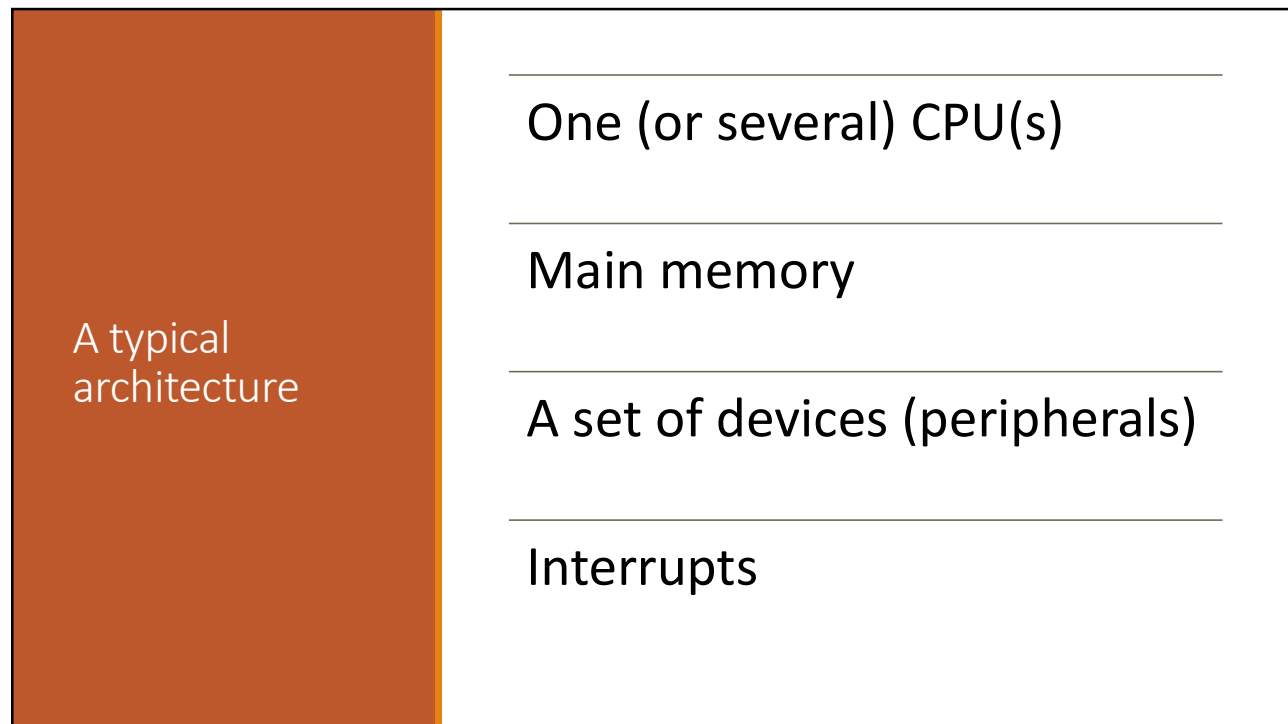




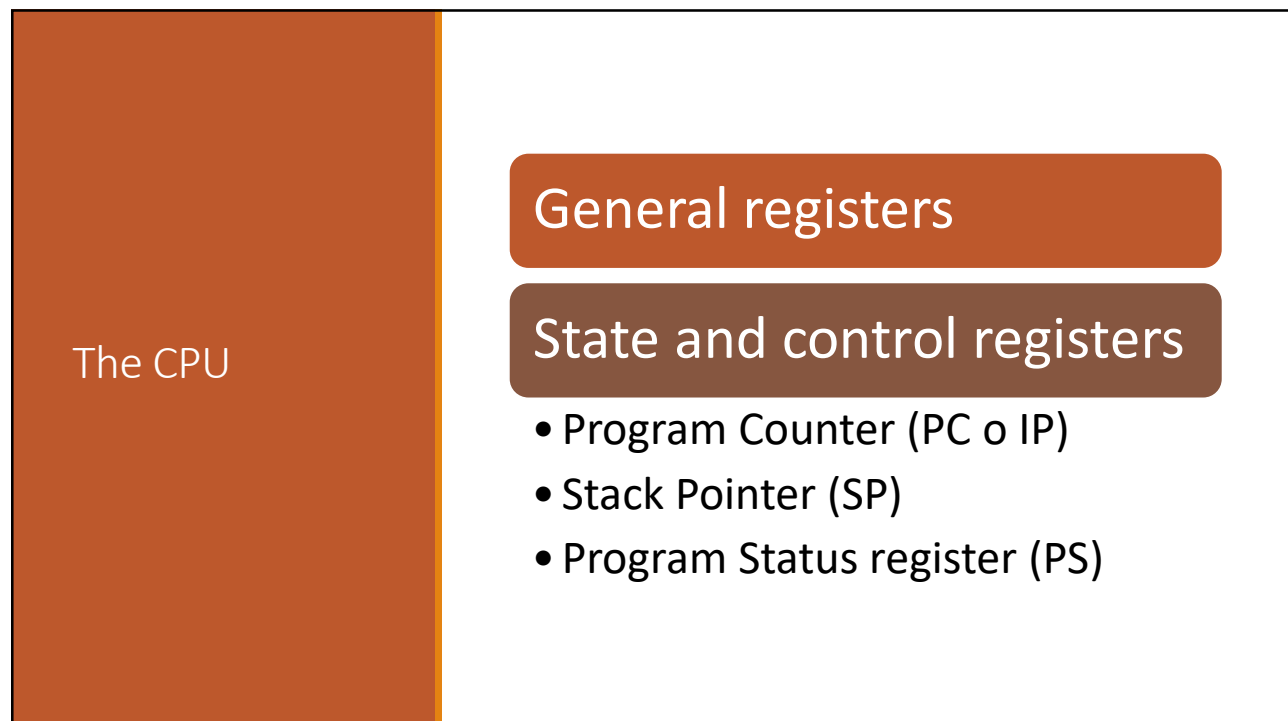
1



2



3



4

The program status register



Condition code

Bit mask that keeps the status of the processor



CPU mode

User mode VS kernel mode



Interrupt enable bit

5

Fetch-execution cycle

If there are pending interrupts and the interrupts are enabled

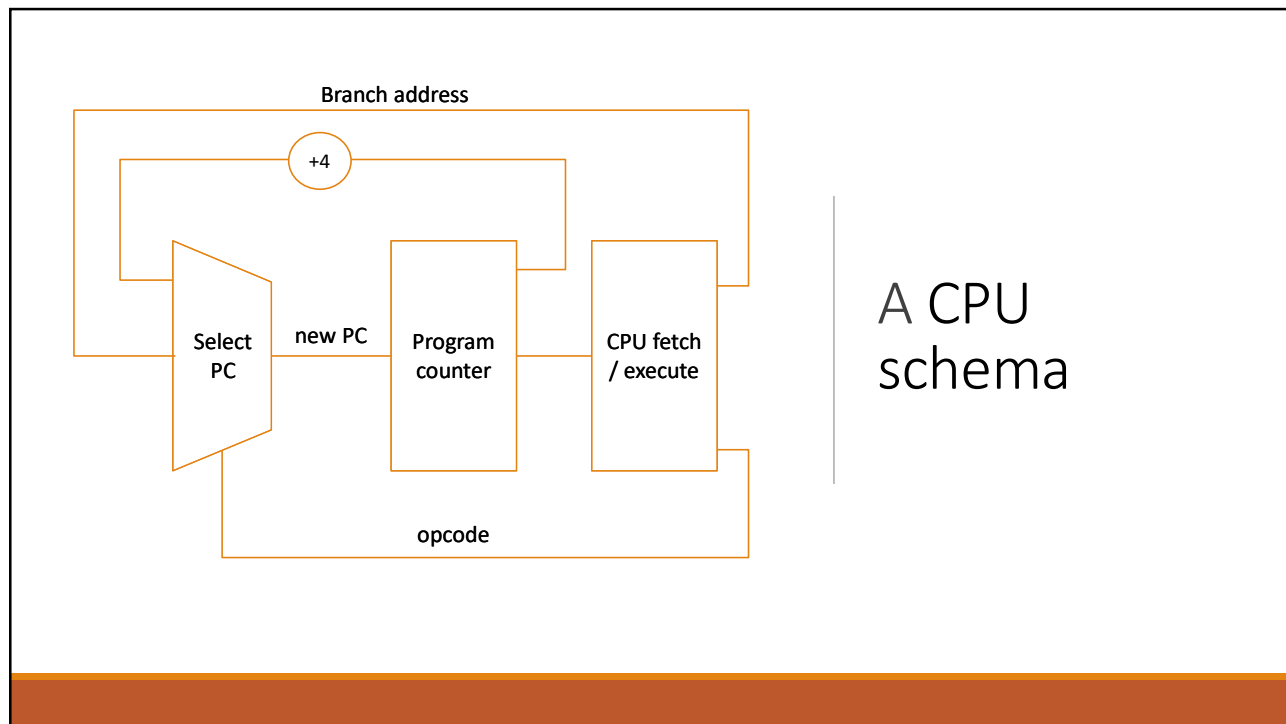
- Manages the interrupt

Else

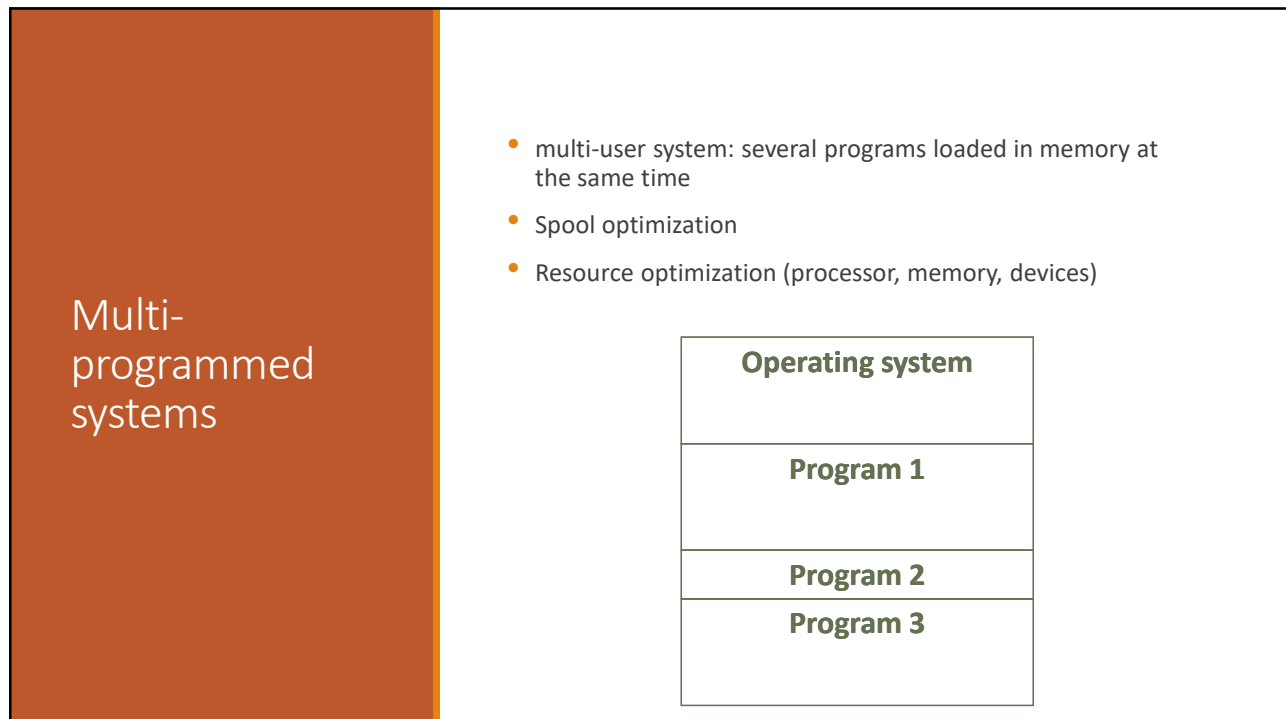
- Loads the instruction at address PC
- Executes the instruction
- $PC = PC + 4$ (*)

(*) assumes that the instruction occupies 4 bytes

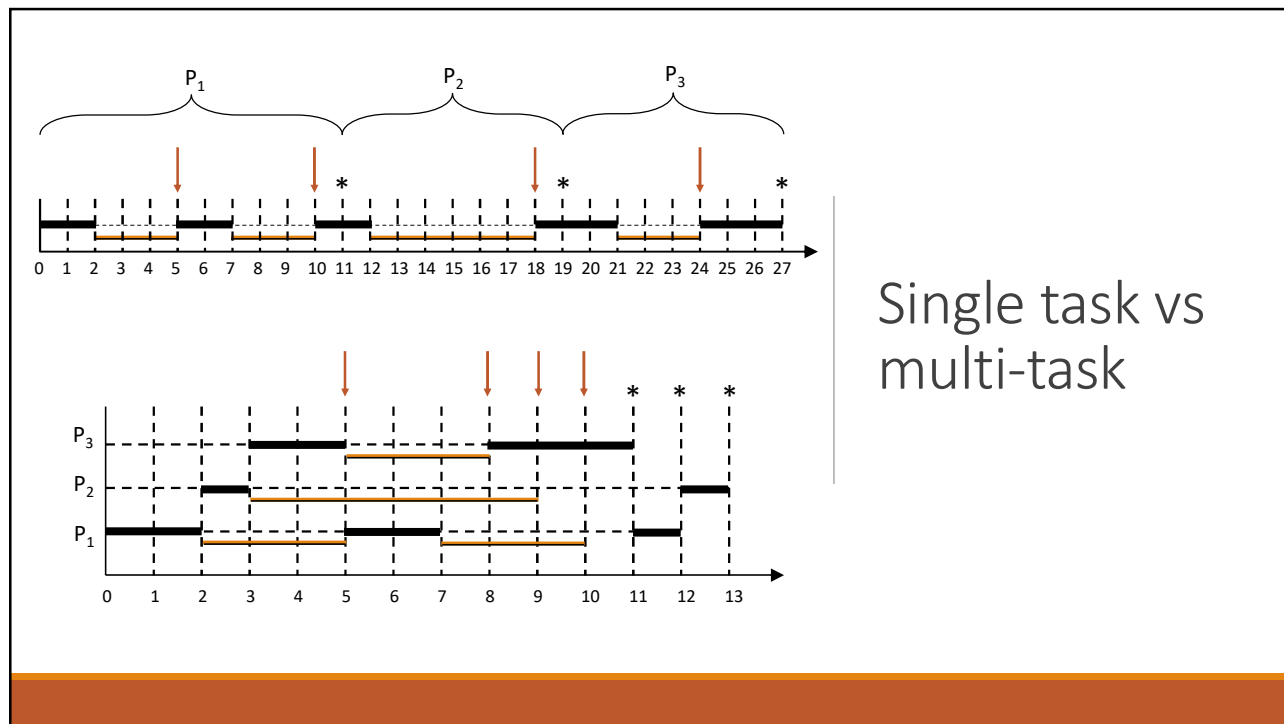
6



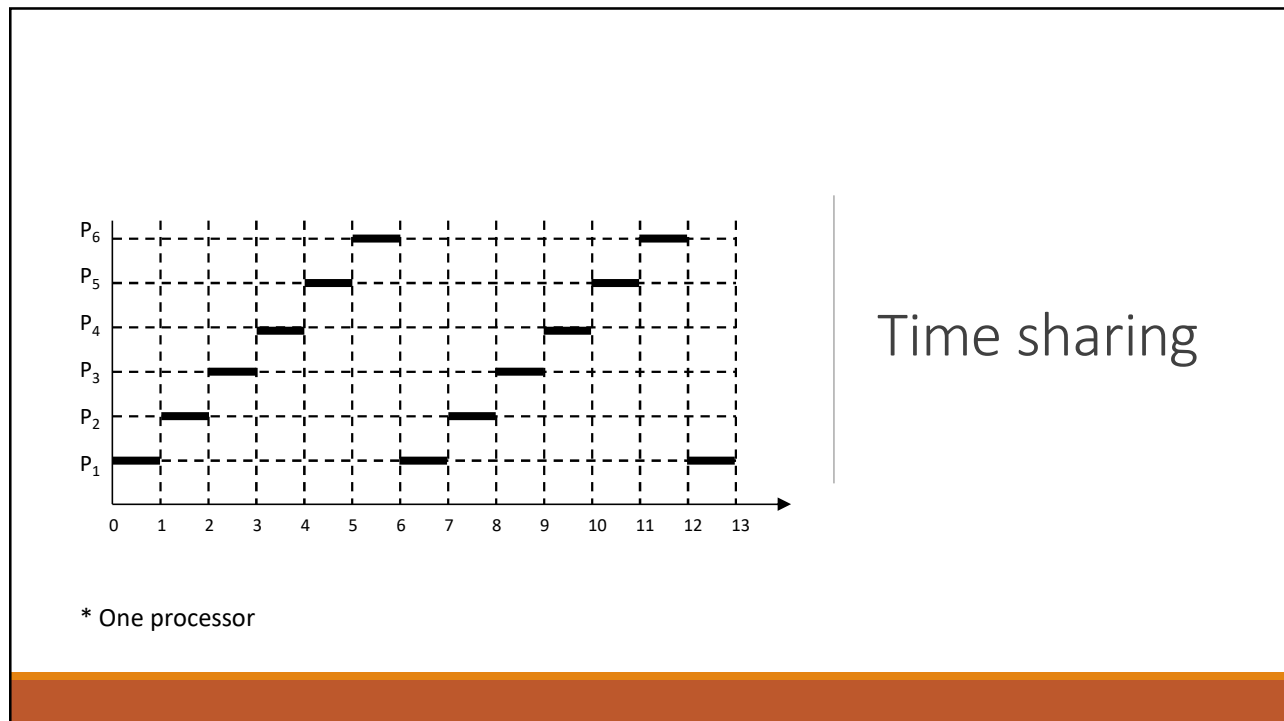
7



8

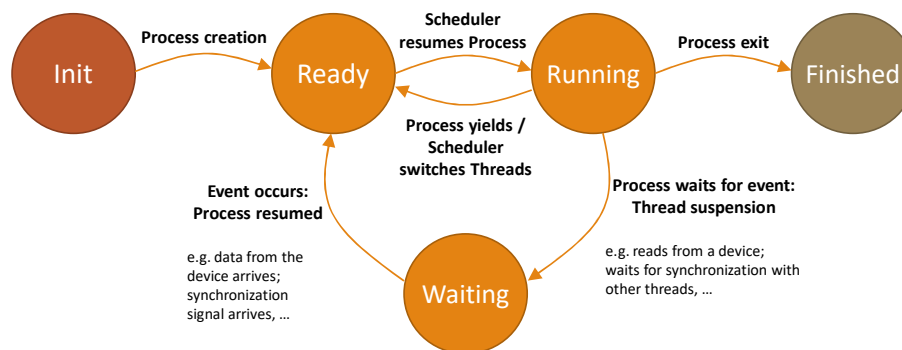


9



10

Processes Lifecycle



11

The Kernel Abstraction

12

Challenge: Protection

How do we execute code with restricted privileges?

- Either because the code is buggy or if it might be malicious

Some examples:

- A script running in a web browser
- A program you just downloaded off the Internet
- A program you just wrote that you haven't tested yet

13

Thought Experiment

Implementing execution with limited privileges:

- Execute each program instruction in a simulator
- If the instruction is permitted, do the instruction
- Otherwise, stop the process
- Basic model in Javascript, ...

How do we go faster?

- Run the unprivileged code directly on the CPU?

14

Main Points



Process concept

A process is an OS abstraction for executing a program with limited privileges



Dual-mode operation: user vs. kernel

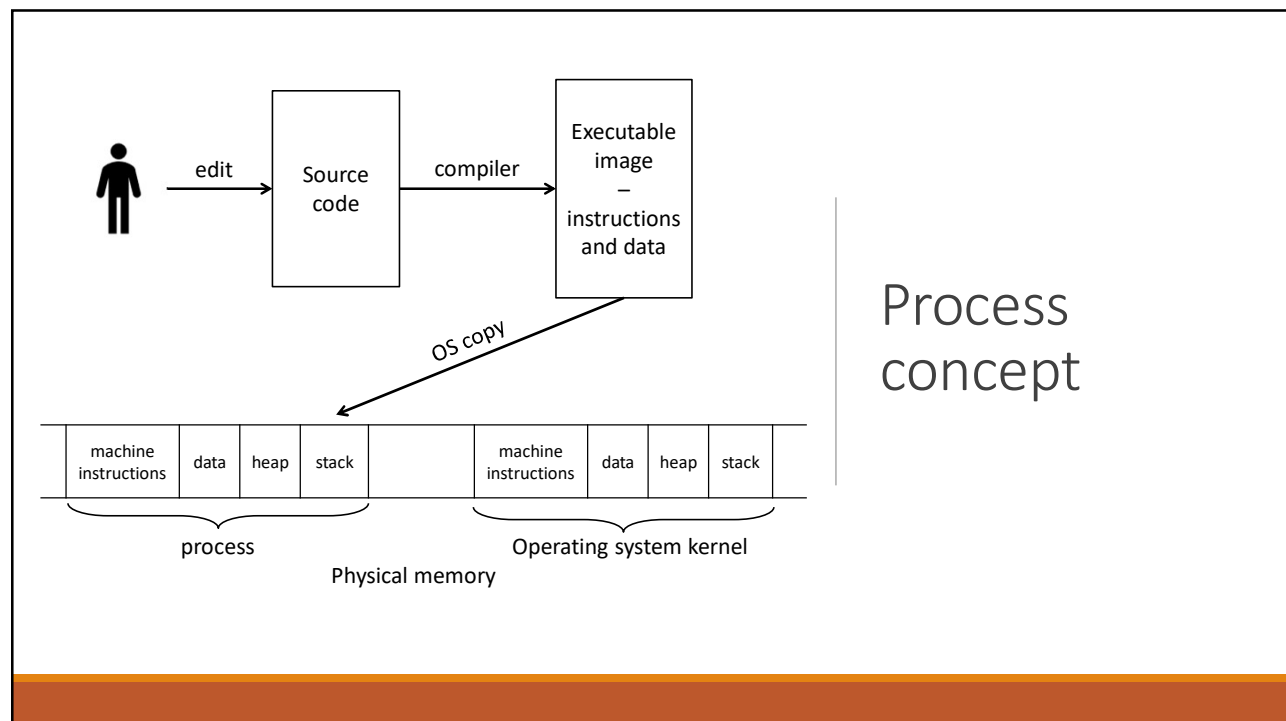
Kernel-mode: execute with complete privileges
User-mode: execute with fewer privileges



Safe control transfer

How do we switch from one mode to the other?

15



16

Process Concept

Process: a sequence of activities activated by a program, **running with limited rights**

- Process control block (PCB): the data structure the OS uses to keep track of a process
- Process Table: contains all PCBs
- Two elements:
 - Thread: executes a sequence of instructions within a process
 - Potentially many threads per process (for now 1:1)
 - Thread aka lightweight process
 - Address space: set of rights of a process
 - Memory that the process can access
 - Other permissions the process has (e.g., which procedure calls it can make, what files it can access)

17

Program and Process

Program: a static sequence of instructions

Process:

- a dynamic entity: a sequence of activities described by a program
- executed on a set of CPUs with limited rights

Several processes can be activated on the same program

- The processes execute the same code
- Each process executes the program on different data and/or in different times

18

Process Control Block

It is a data structure:

Process name

- Can be the index of the PCB in the process table

Pointers to process threads

Assigned memory

Other assigned resources

- Files, devices, etc...

19

Hardware Support: Dual-Mode Operation

Kernel mode

- Execution with the full privileges of the hardware
- Read/write to any memory, access any I/O device, read/write any disk sector, send/read any packet

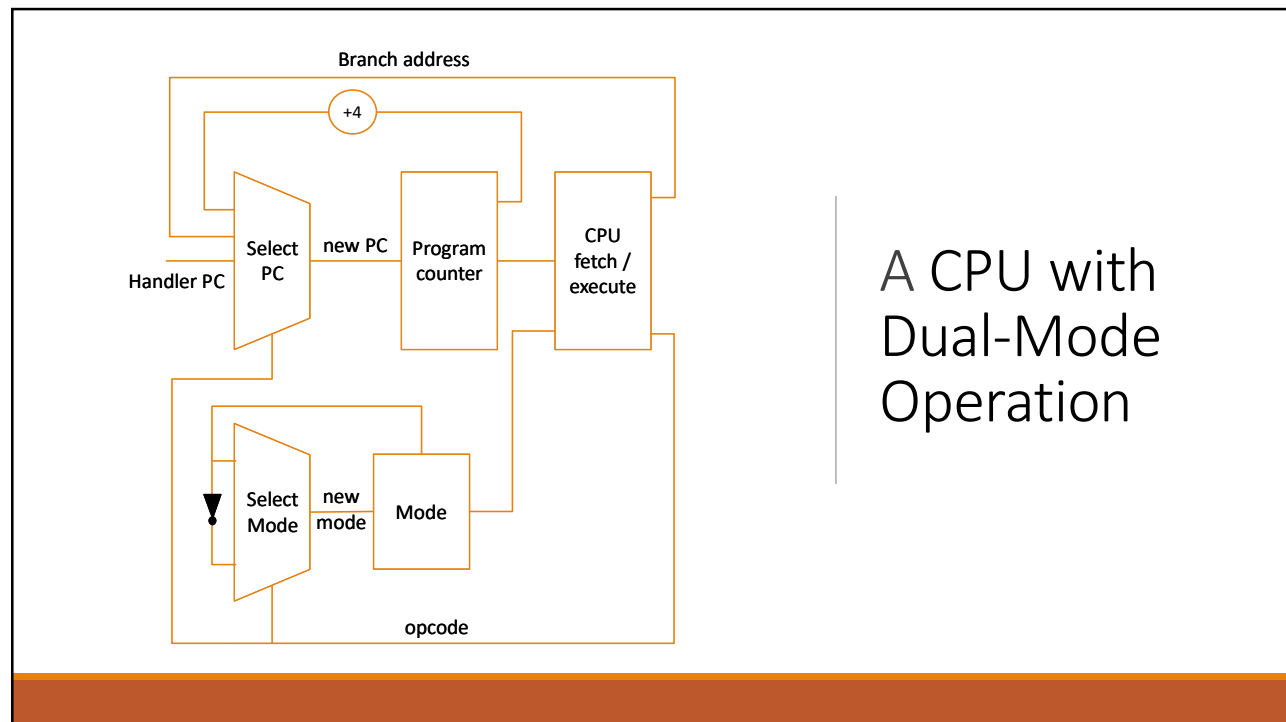
User mode

- Limited privileges
- Only those granted by the operating system kernel

On the x86, mode stored in EFLAGS register

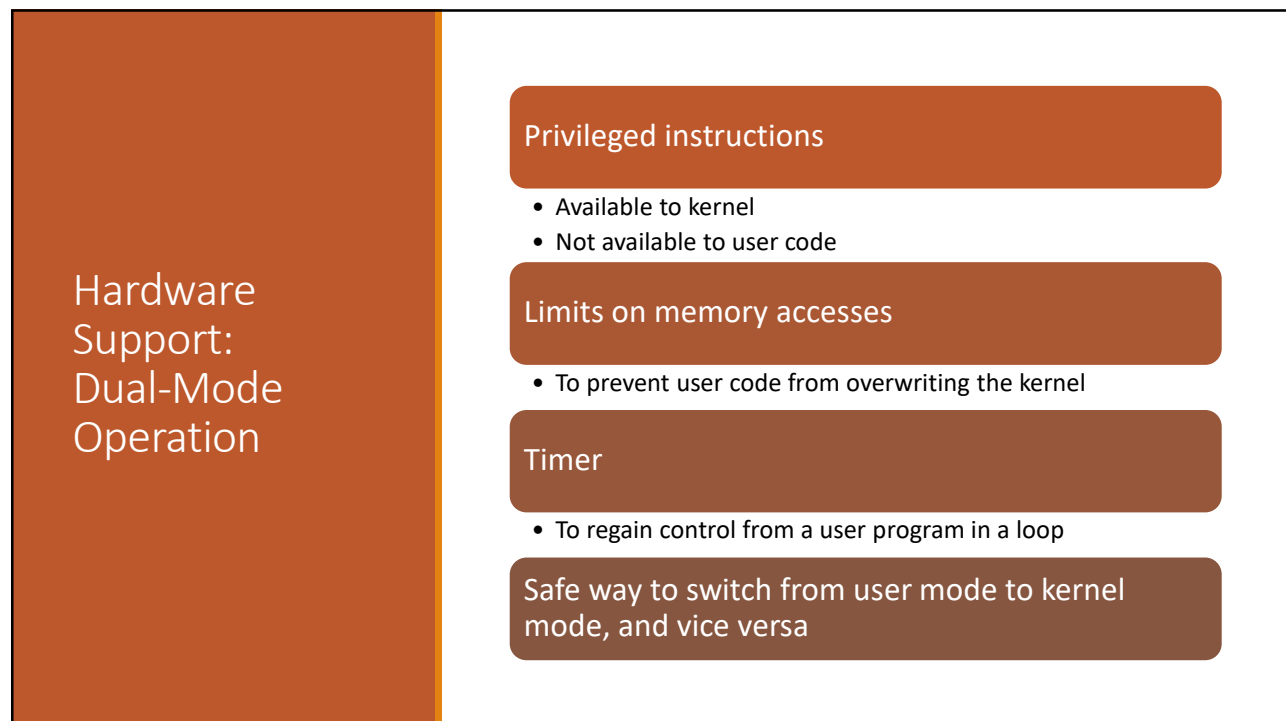
- In general in the Program Status Register

20

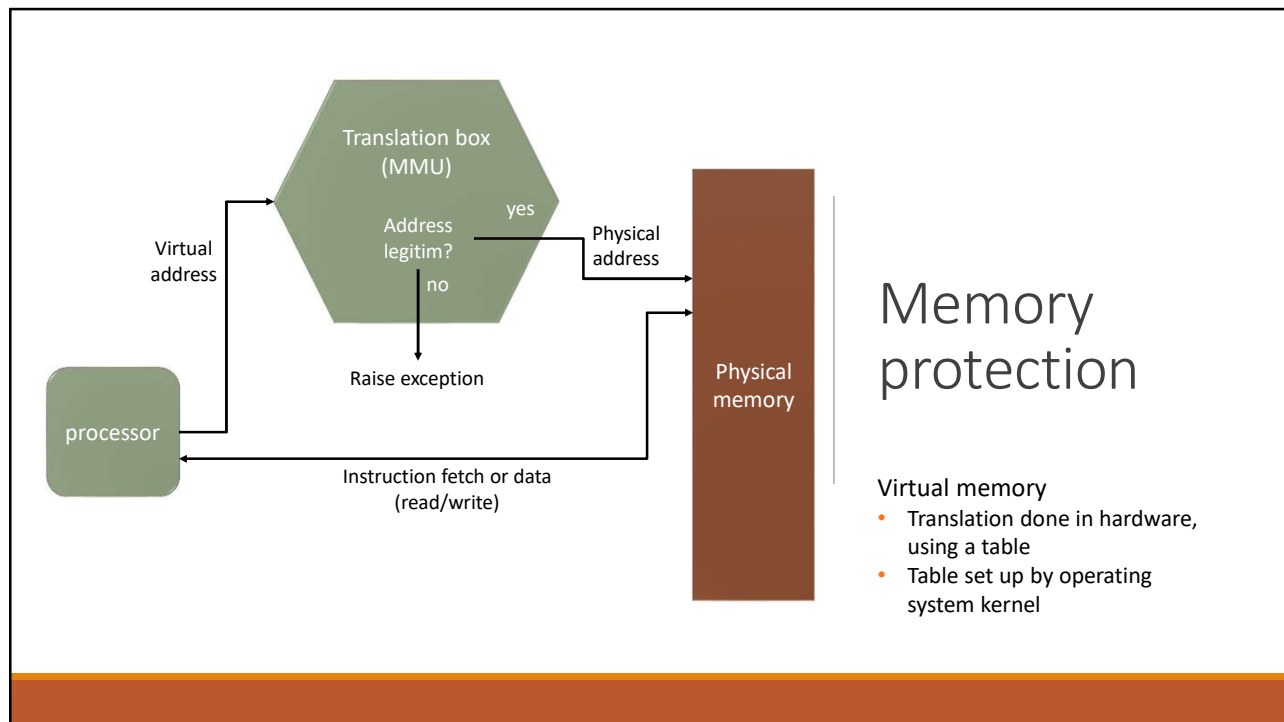


A CPU with Dual-Mode Operation

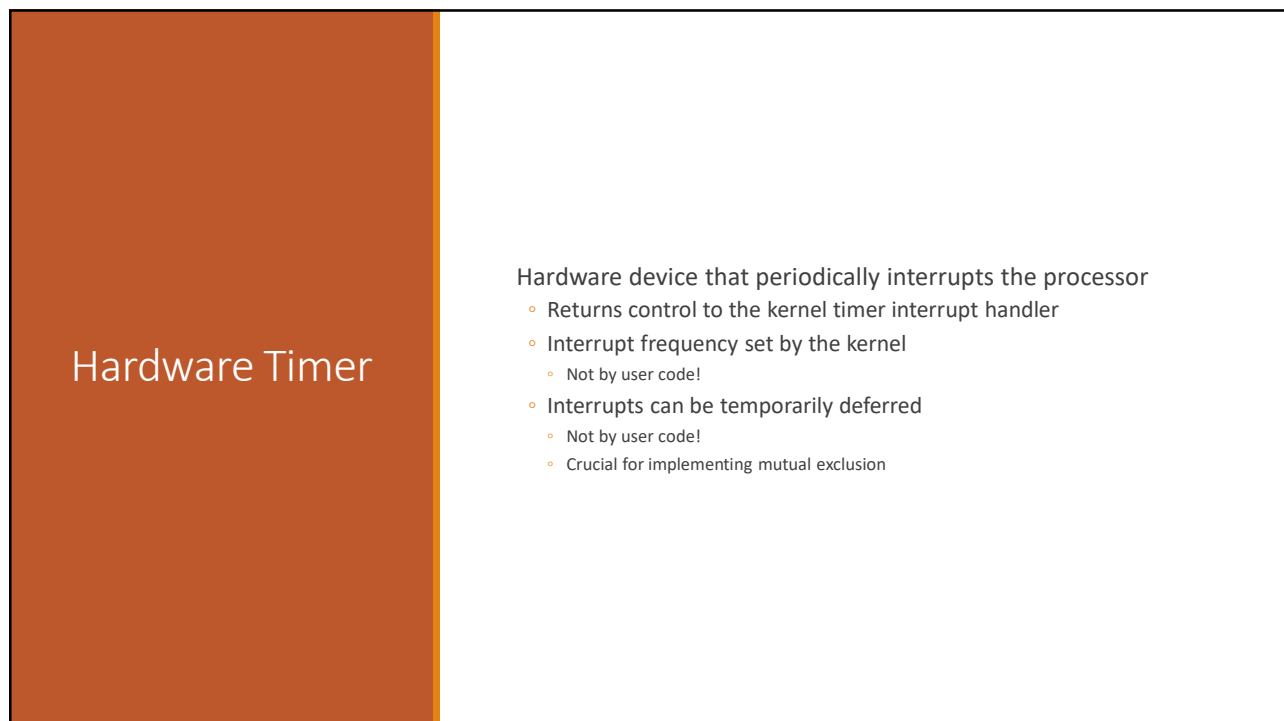
21



22



23



24

Mode Switch

From user-mode to kernel

- Interrupts
 - Triggered by timer and I/O devices
- Exceptions
 - Triggered by unexpected program behavior
 - Or malicious behavior!
- System calls (aka protected procedure call)
 - Request by program for kernel to do some operation on its behalf
 - Only limited # of very carefully coded entry points

25

Mode Switch

From kernel-mode to user-mode

- Return from interrupt, exception, system call
 - Resume suspended execution
- New process/new thread start
 - Jump to first instruction in program/thread
- Process/thread context switch
 - Resume some other process
- User-level upcall
 - Asynchronous notification to user program

26

How do we take interrupts safely?

Interrupt vector

- Limited number of entry points into kernel

Kernel interrupt stack

- Handler works regardless of state of user code

Interrupt masking

- Handler is non-blocking

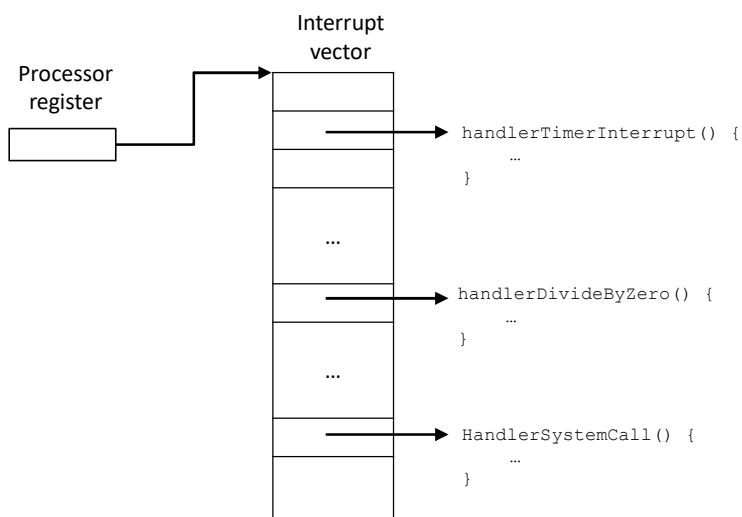
Atomic transfer of control

- Single instruction to change:
 - Program counter
 - Stack pointer
 - Memory protection
 - Kernel/user mode

Transparent restartable execution

- User program does not know interrupt occurred

27



Interrupt vector

Table set up by OS kernel; pointers to code to run on different events

28

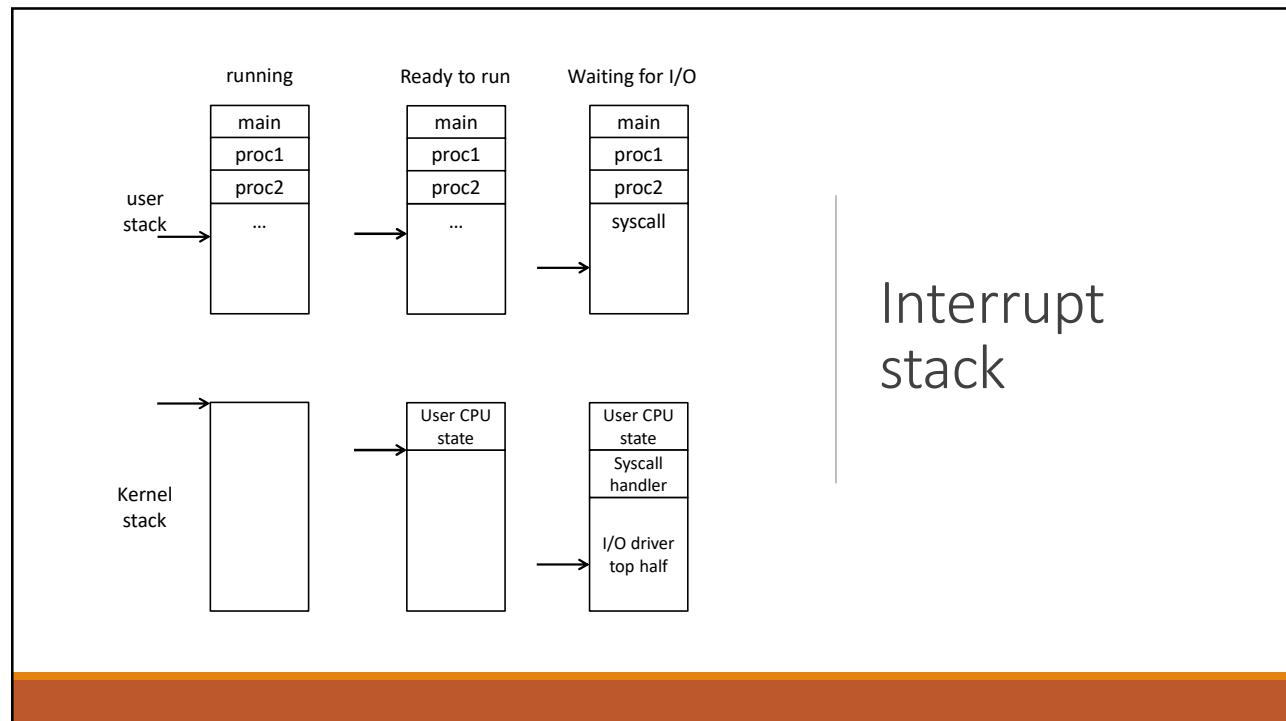
Interrupt Stack

Per-processor, located in kernel (not user) memory

- Usually a thread has both: kernel and user stack

Why can't interrupt handler run on the stack of the interrupted user process?

29



30

Interrupt Masking

Interrupt handler runs with interrupts off

- Re-enabled when interrupt completes

OS kernel can also turn interrupts off

- Eg., when determining the next process/thread to run
- If defer interrupts too long, can drop I/O events
- On x86
 - CLI: disable interrupts
 - STI: enable interrupts
 - Only applies to the current CPU

Cf. implementing synchronization, chapter 5

31

Interrupt Handlers

Non-blocking, run to completion

- Minimum necessary to allow device to take next interrupt
- Any waiting must be of limited duration
- Wake up other threads to do any real work

Rest of device driver runs as a kernel thread

- Queues work for interrupt handler
- (Sometimes) wait for interrupt to occur

32

Atomic Mode Transfer

On interrupt (x86):

- Save current stack pointer
- Save current program counter
- Save current processor status word (condition codes)
- Switch to kernel mode
- Vector through interrupt table
- Switch to handler PC & kernel PSW
- Interrupt handler saves registers it might clobber

In Hardware

33

Before

User-level process

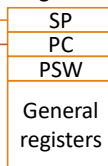
Code:

```
Foo() {
  while(...) {
    x=x+1;
    y=y-2;
  }
}
```

Stack:



Registers:



Kernel

Code:

```
handler() {
  push A
  ...
}
```

Exception Stack:



34

During

User-level
process

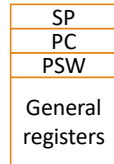
Code:

```
Foo() {
  while(...) {
    x=x+1;
    y=y-2;
  }
}
```

Stack:



Registers:

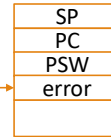


Kernel

Code:

```
handler() {
  push A
  ...
}
```

Exception
Stack:



35

After

User-level
process

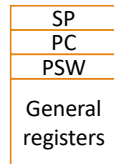
Code:

```
Foo() {
  while(...) {
    x=x+1;
    y=y-2;
  }
}
```

Stack:



Registers:

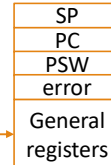


Kernel

Code:

```
handler() {
  push A
  ...
}
```

Exception
Stack:



36

At end of
handler

Handler restores saved registers

Atomically return to interrupted process/ thread (IRET instruction)

- Restore program counter
- Restore program stack
- Restore processor status word/condition codes
- Switch to user mode
- Enable interrupts

37

Interrupt management a simple example

Initial state: interrupt '500' occurs when executing instruction A000

Registri nella CPU

PC	A000
PS	PSW P
SP	FFFF
R1	AAAA
R2	BBBB
...	

Memoria

programma P

....	...
A000	istr. 1
A004	istr. 2
A008	istr. 3
A016	istr. 4
A020	istr. 5
...	...

Stack di P

...	...
FFF0	
FFF3	
FFF7	
FFFB	
FFFF	...

stack nel nucleo

...	...
2996	
2997	
2998	
2999	
3000	

Interrupt Handler

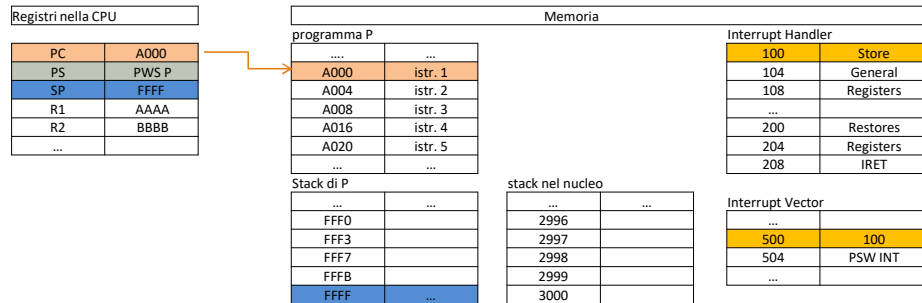
100	Store
104	General
108	Registers
...	
200	Restores
204	Registers
208	IRET

Interrupt Vector

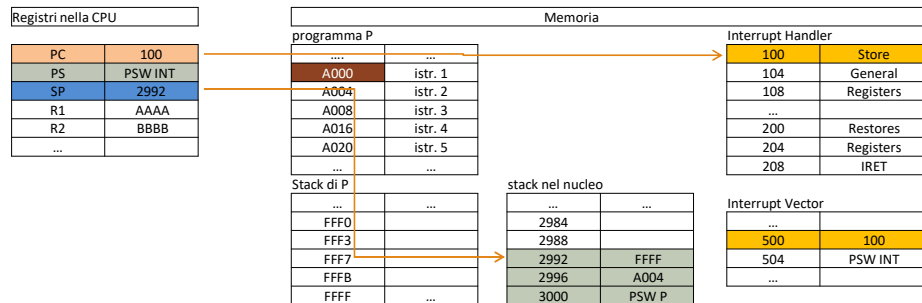
...	
500	100
504	PSW INT
...	

38

1) Initial state

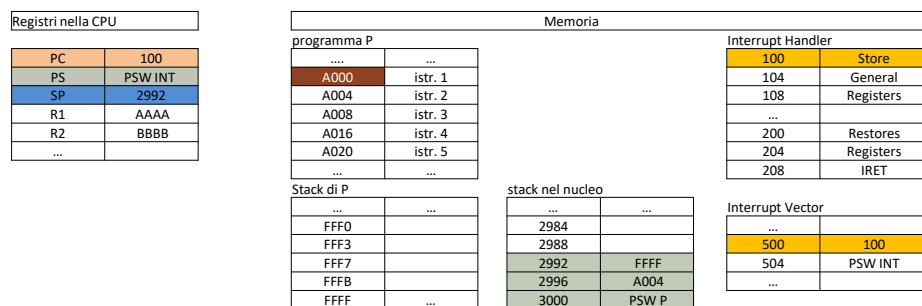


2) Interrupt recognized after instruction A000 (PC incremented to A004)

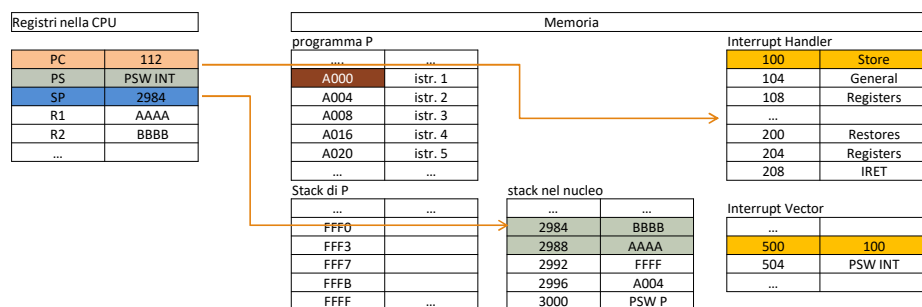


39

2) Interrupt recognized after instruction A000

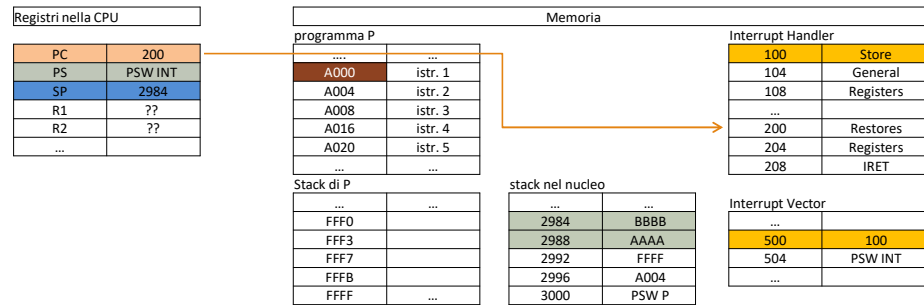


3) Stores general registers

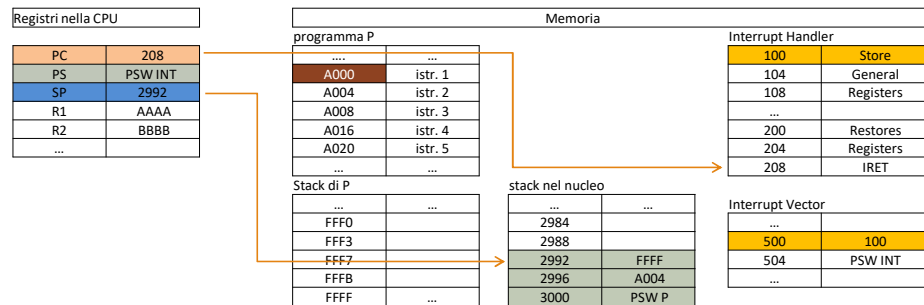


40

4) Executes interrupt handler

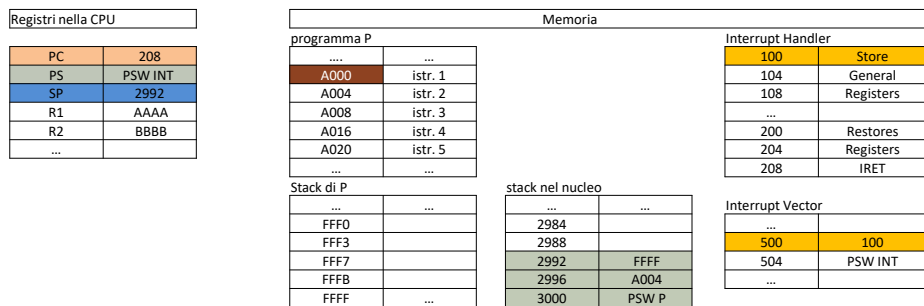


5) Restores general registers

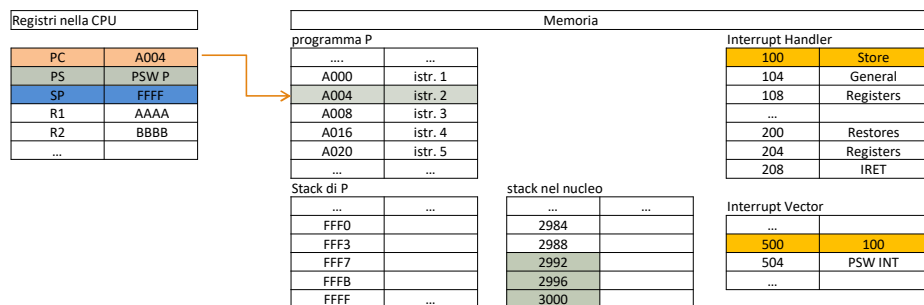


41

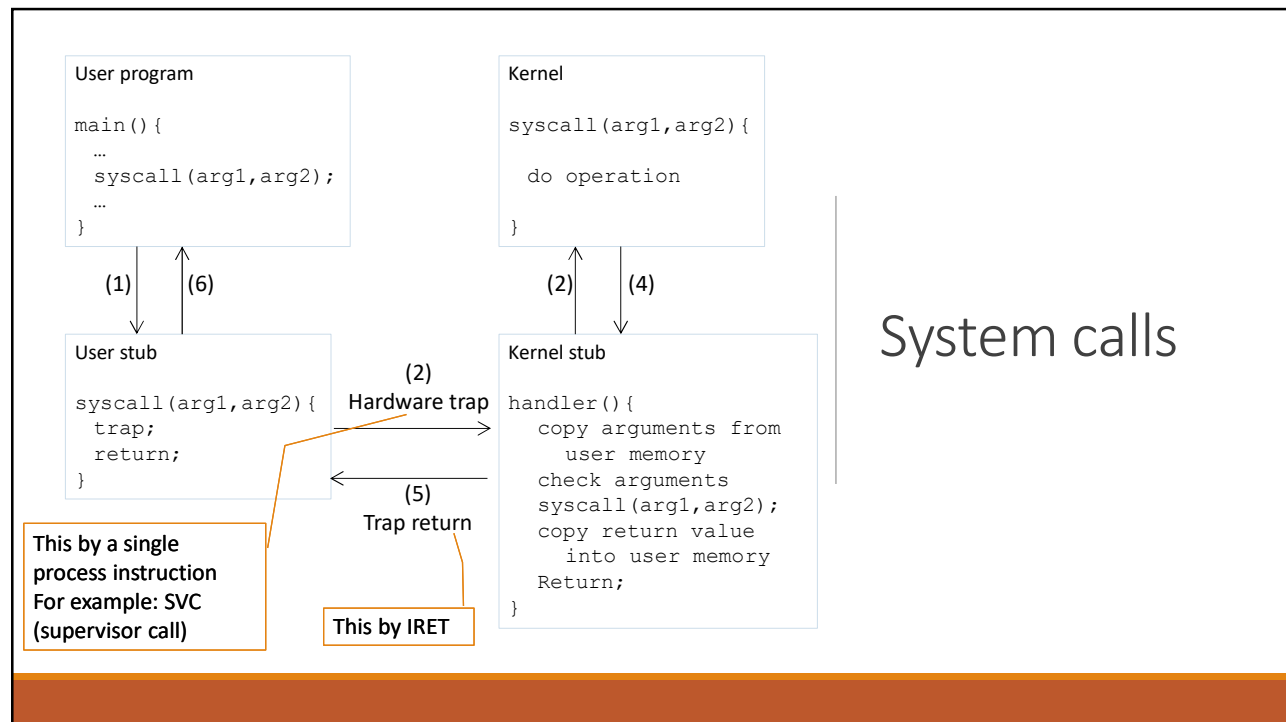
5) Restores general registers



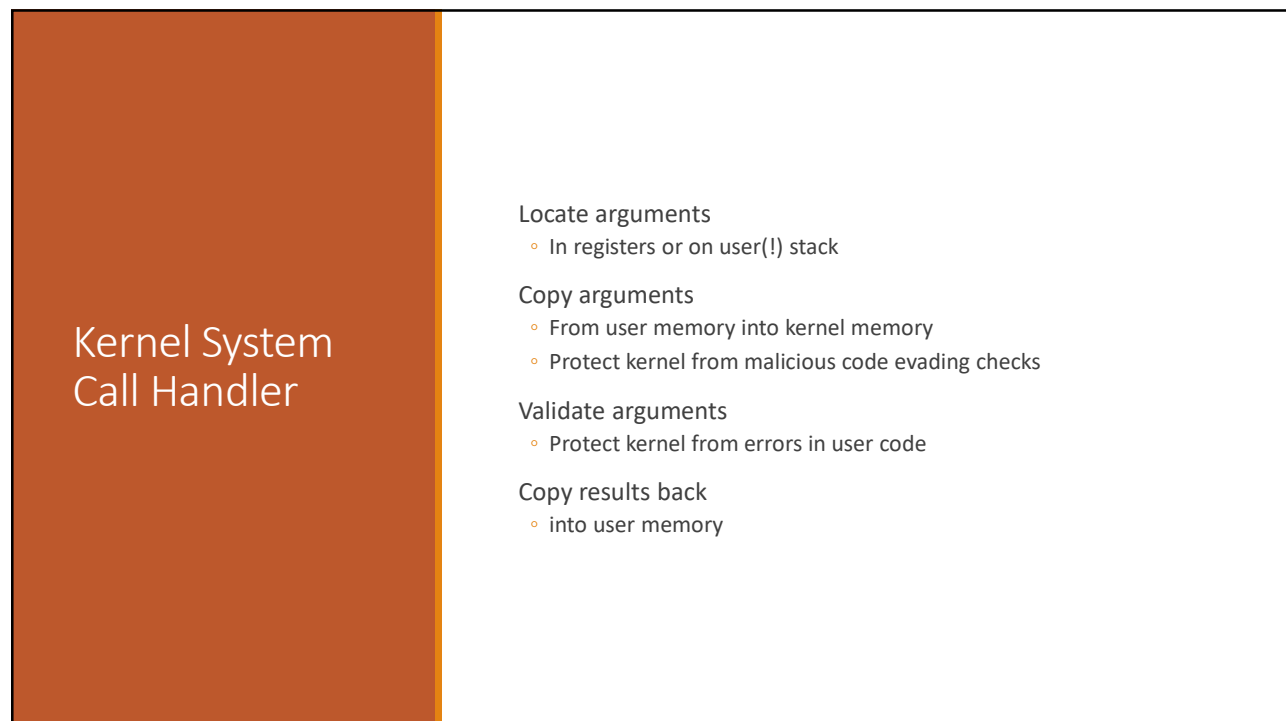
6) Executes IRET



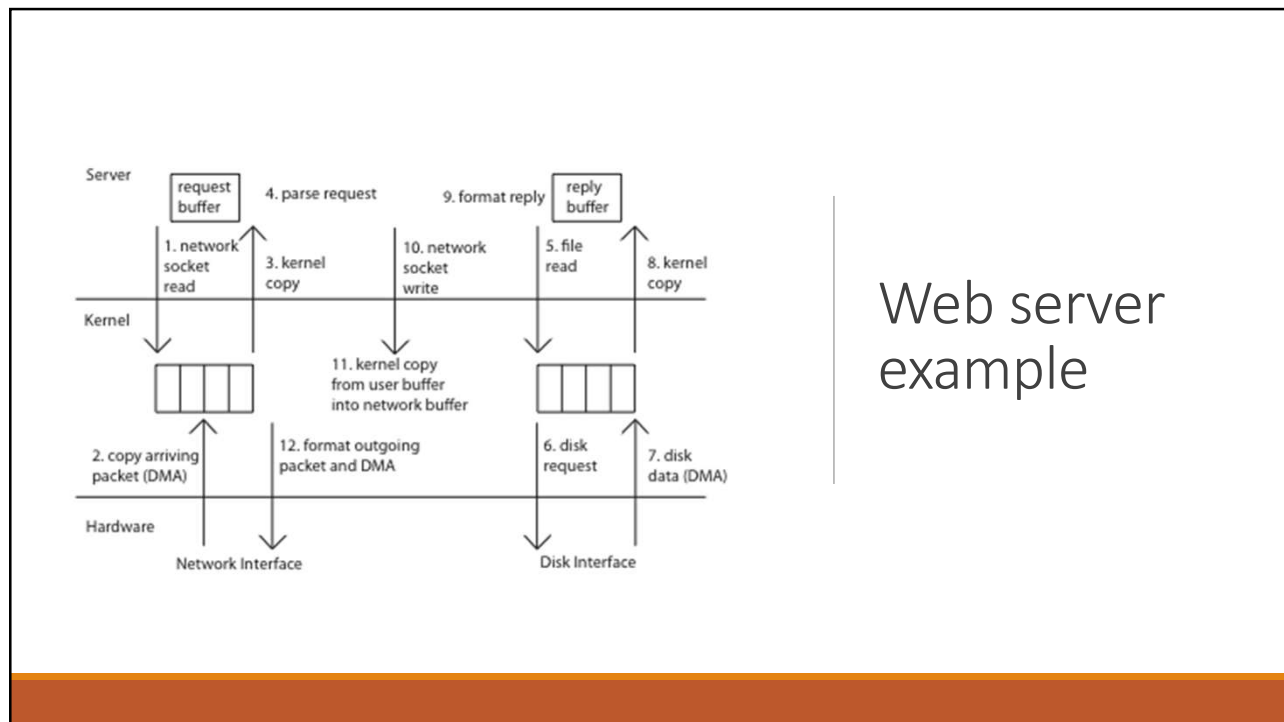
42



43



44



Web server example

45

Upcall: User-level interrupt

Also known as UNIX signals

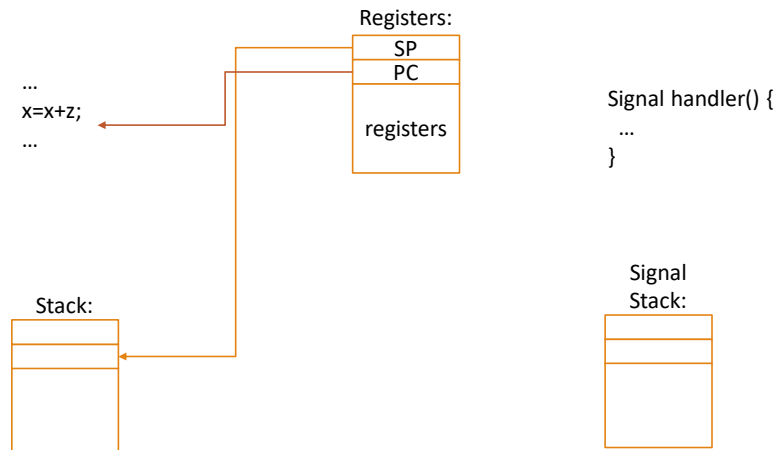
- Notify user process of event that needs to be handled right away
- Time-slice for user-level thread manager
- Interrupt delivery for VM player (see later)

Direct analogue of kernel interrupts

- Signal handlers – fixed entry points
- Separate signal stack
- Automatic save/restore registers – transparent resume
- Signal masking: signals disabled while in signal handler

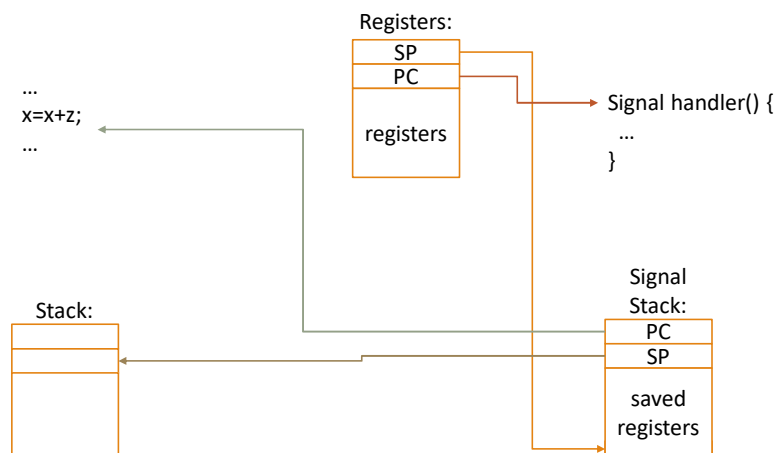
46

Upcall: before



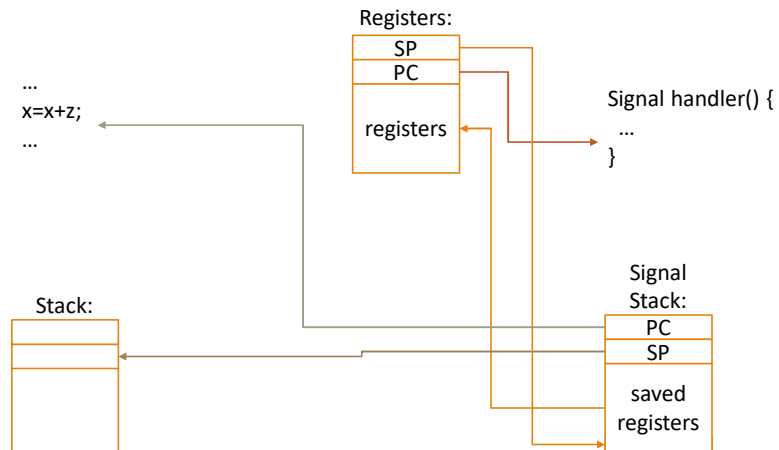
47

Upcall: after

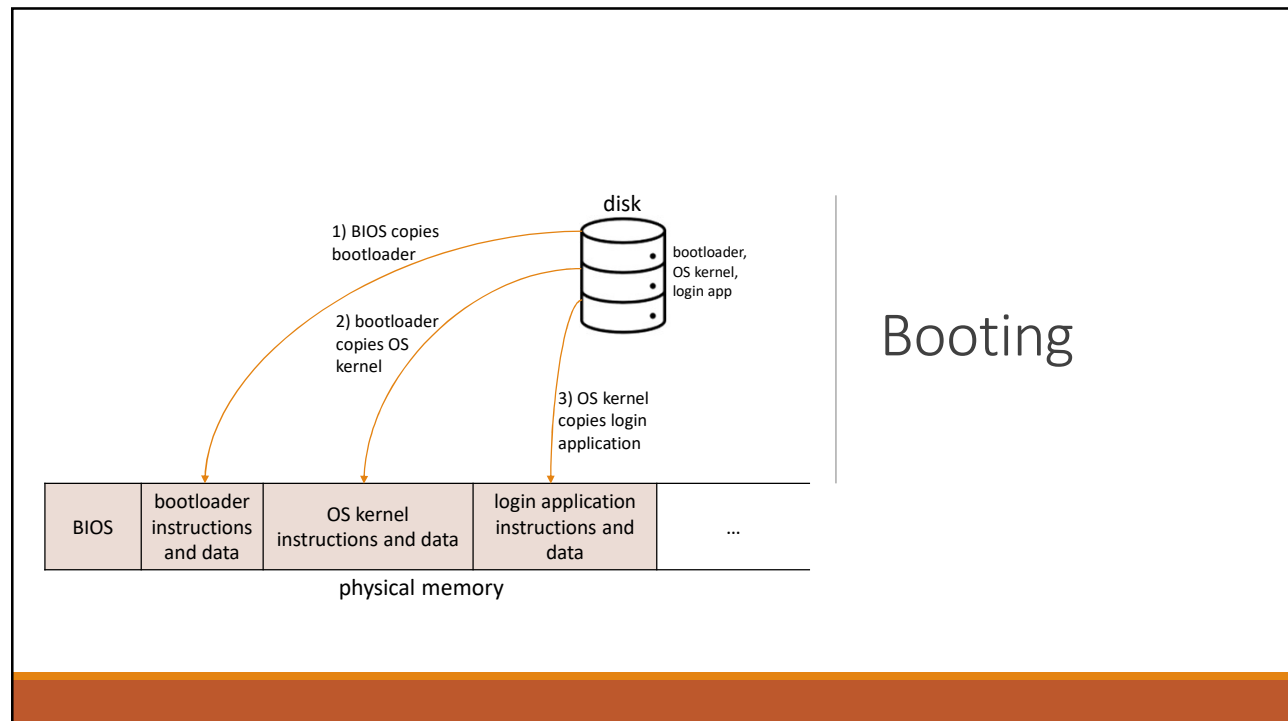


48

Upcall: end of handler



49



50

Notes on FAT file system

51

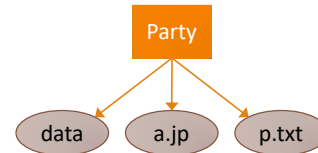
Named Data in a FAT File System



52

Directories

- Each file (and directory) belongs to a directory;
- Each directory is a **data structure** that links file names to file attributes
 - Example of attributes: file size, address on disk, access rights, time of last access, time of creation, ...



A possible implementation (**FAT file systems**):

- The directory is a table, it associates each file name to its file descriptor (which includes all attributes of the file)

name	descriptor
data	Attribute: type, first data block, etc.
a.jp	Attribute: type, first data block, etc.
p.txt	Attribute: type, first data block, etc.

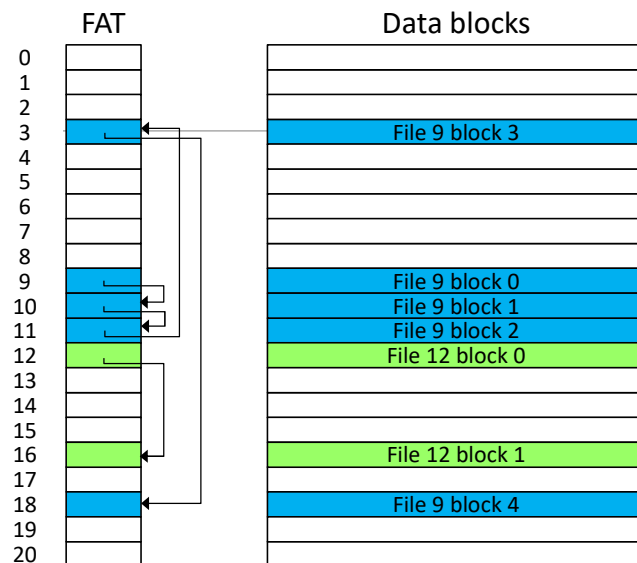
Implementation of directory Party

53

FAT file system

File Allocation Table (FAT) :
 Linked list index structure
 Simple, easy to implement
 Still widely used (e.g., thumb drives)

File table:
 Linear map of all blocks on disk
 Each file a linked list of blocks



54

Physical disk organization with FAT

