



# **S E C U R I T Y P O L I C I E S**

AND THEIR  
ENFORCEMENTS

# Secure Programming Languages

## Stack, Heap and code area

- Stack Canaries
- Reorder layout of AR and Heap elements (randomization)
- Shadow stack
- NX memory
- FAT Pointer & Data Descriptor

## Static Analysis

- Monotone Framework

## ○ What we have learned (example)

- Stack canaries:
  - **Run-time support**
  - **Code instrumentation checking at run-time the compliance with the security policy**
- **Stack Canary Security policy: the control flow cannot be altered**



# ○ A motivating example

## A TRUSTED PROGRAM

```
:
```

```
10: z := Mem(x)
```

```
11: if(i ≥ 0) jump y
```

```
:
```

*The program dereferences memory  
and makes use of indirect control flow transfer.*

The attacker sets things up in a way so that  
unauthorized memory is copied into x

the attacker can later access

## THE ATTACKER CODE

```
20: i:=0
```

```
21: x := attacker's desired address
```

```
22: y:=24
```

```
23: f(0 = 0) jump 10
```

```
24: copy memory contents from z
```



## ○ A motivating example

### A TRUSTED PROGRAM

```
:
```

```
10: z := Mem(x)
```

```
11: if(i ≥ 0) jump y
```

```
:
```

### THE ATTACKER CODE

```
20: i:=0
```

```
21: x := attacker's desired address
```

```
22: y:=24
```

```
23: f(0 = 0) jump 10
```

```
24: copy memory contents from z
```

*Form of return-oriented programming (ROP): the attacker knows program text.*

*This relies on the ability to change control flow using indirection so that commands are executed in the order needed by the attackers to carry out their goals.*

## ○ Countermeasure: sandbox

We (aka the compiler) instruments the original program with “code sandbox” to manage flow control between commands at program point  $l$  and program point  $h$  (i.e.  $pc_l$  and  $pc_h$ ).

### **The instrumentation**

The compiler rewrites indirect jump commands to ensure that their target always lies within the right bounds ( $l, h$ ).


*Rewrite all  $\text{if}(Q) \text{ jump } e$  commands as  $\text{if}(Q) \text{ jump } (e \ \& \ pc_h) \mid pc_l$*



# ○ Our approach



Enforce a general safety property that takes aspects of control flow and program state into account.



Example, suppose that our programming language has the ability to make three types of function calls,

**send and receive from the network**  
**read from a local file.**



Then we want to enforce a safety policy on untrusted code which says that send cannot be called after read.






## Next step: security policies

Introduce suitable mechanisms for

1) Defining security policies

2) Implementing enforcement mechanisms via compiler instrumentation





# ○ Some challenging questions

Can we prove that mechanism **M** enforces the security policy **P**?

- What is the mathematical definition of a policy?
- What is the programming abstractions for declaring security policies?
- What does it mean to enforce a policy within a programming language?

Are there limits to what is enforceable?

- Which enforcement approaches are best suited to which

Policies?

- Are there some policies that are completely beyond any known enforcement strategy?
- Are some enforcement approaches strictly more powerful than others?

OVERALL

- what is the landscape of policies, policy classes, and enforcement mechanisms?





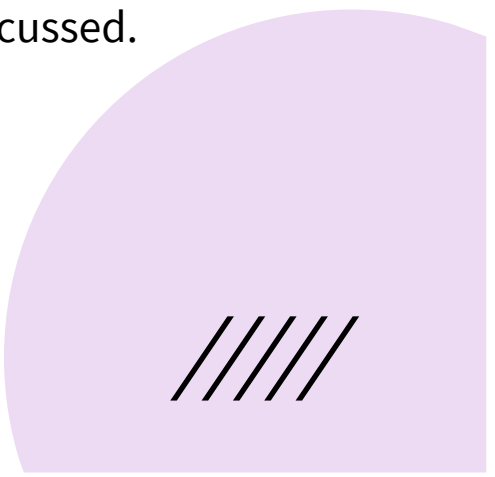
Starting point


## Enforceable Security Policies

F. Schneider, [ACM Transactions on Information and System Security](#), February 2000

### Abstract

A precise characterization is given for the class of security policies enforceable with mechanisms that work by monitoring system execution, and automata are introduced for specifying exactly that class of security policies. Techniques to enforce security policies specified by such automata are also discussed.





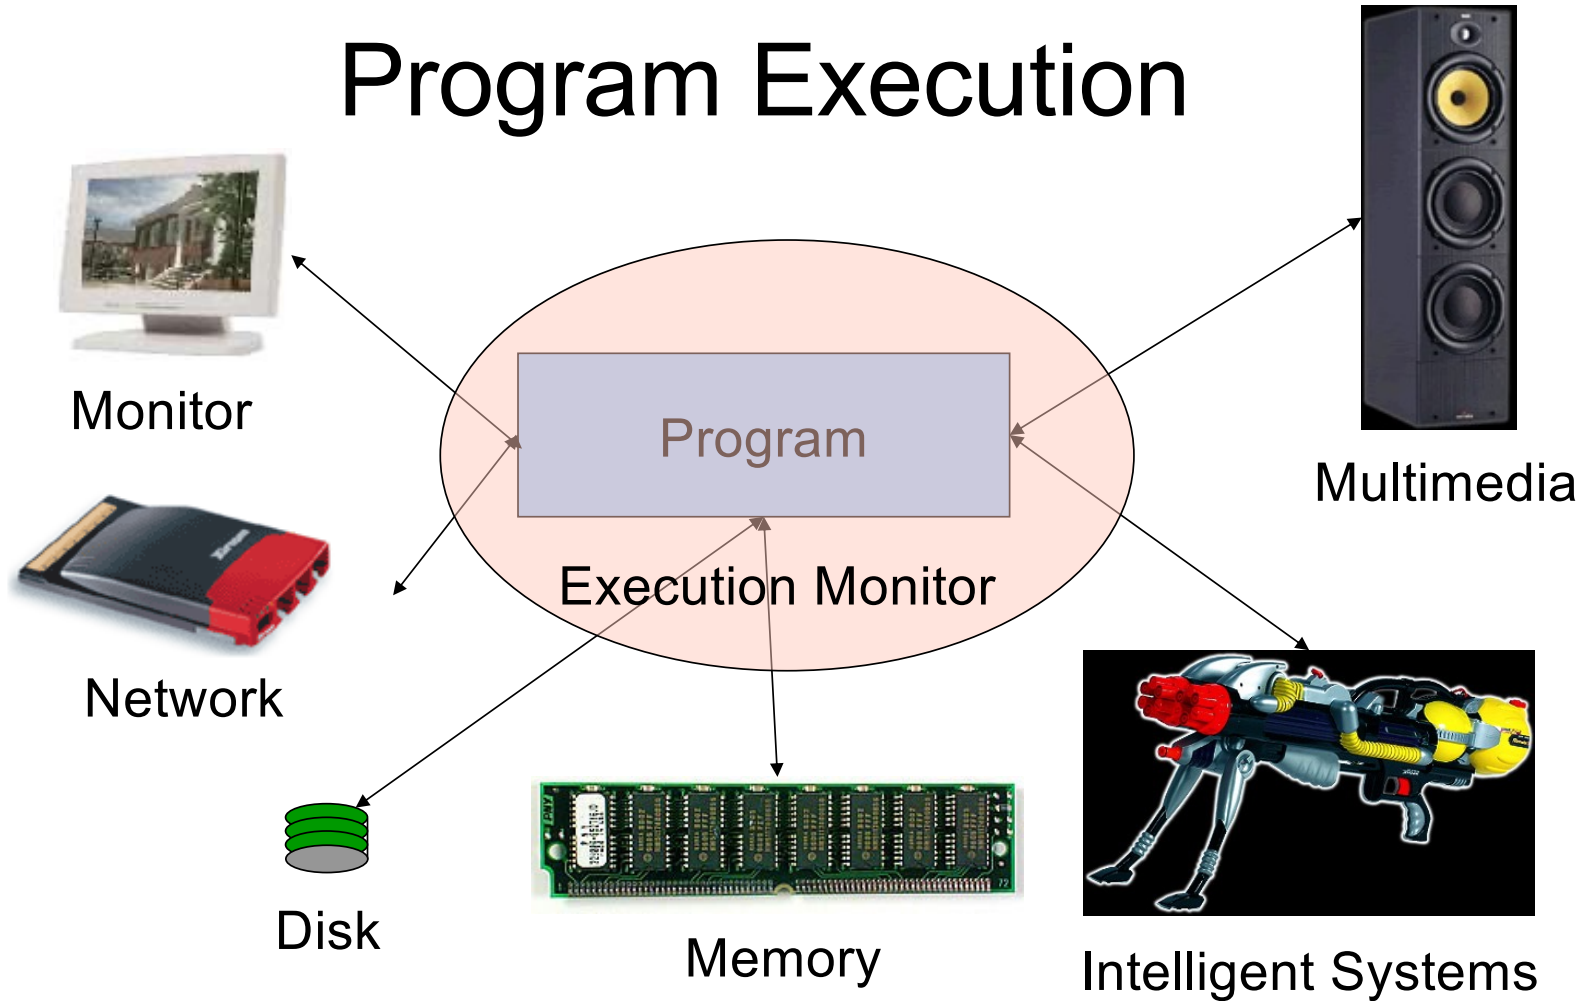
### Execution Monitors (EMs)

- EMs watch untrusted programs at runtime
- Events (raised by the run time executions) are mediated by the EM
- Violations solicit EM interventions (e.g. termination)

### Example: File system access control

- EM is inside the OS
- decides policy violations using access control lists (ACLs)

# The Context: Program Execution



# Ideal Execution Monitor

1. **Sees *everything* a program is about to do before it does it**
2. **Can *instantly* and *completely* stop program execution (or prevent action)**
3. **Has *no other effect* on the program or system**

Can we build this?

Probably not .....

# Real ~~Ideal~~ Execution Monitor


**most things**

1. Sees ~~everything~~ a program is about to do before it does it
2. Can ~~instantly and completely stop~~ program execution (or prevent action)
3. Has ~~no other~~ **limited** effect on the program or system

# Operating Systems

- Provide execution monitors for most security-critical resources
  - When a program opens a file in Unix or Windows, the OS checks that the principal running the program can open that file
- Doesn't allow different policies for different programs
- No flexibility over what is monitored
  - OS decides for everyone
  - Hence, can't monitor inexpensive operations

- OS Execution Monitor









# Execution Monitor


- EMs are run-time modules that runs in parallel with an application
  - monitors may detect, prevent, and recover from application errors at run time
  - monitor decisions may be based on **execution history**



- EM: Good Operations




- EM: Bad Operation



# ○ Intuition

**program's execution on a given input as a sequence of runtime events**



# ○ What policies can be enforced?

- Assume:
  - Security Automaton can see entire state of world, everything about instruction about to execute
  - Security Automaton has unlimited memory, can do unlimited computation
- Are there interesting policies that still can't be enforced?

# ○ What's a Security Policy?

- What's a program?
  - A set of possible executions
- What's an execution?
  - A sequence of states
- What's a security policy?
  - A predicate on a set of executions

# ○ Programs and Policies

- An *execution (or trace)*  $s$  is a sequence of security-relevant program events  $e$  (also called actions)
- Sequence may be **finite** or **(countably) infinite**
  - $s = e_1; e_2; \dots; e_k; e_{\text{halt}}$
  - $s = e_1; e_2; \dots; e_k; \dots$
  - The empty sequence  $\varepsilon$  is an execution
  - If  $s$  is the execution  $e_1; e_2; \dots; e_i; \dots e_i; \dots$  then  $s[i]$  is the execution  $e_1; e_2; \dots; e_i$
- We simplify the formalism.
  - We model program termination as an infinite repetition of  $e_{\text{halt}}$  event.
  - Result: now all executions are infinite length sequences



# ○ Programs and Policies

- A program **S** is a **set** of sequences (possible executions)
  - A program is modelled as the set  $S = \{s_1, s_2, \dots\}$
- A policy **P** is a **property** of programs
- A policy partitions the program space into two groups:
  - Permissible
  - Impermissible
- Impermissible programs are censored somehow (e.g., terminated on violating runs)








## Formally...


- $\Psi$  : set of all possible executions (can be infinite)
- $\Sigma_S$ : set of executions possible by target program  $S$
- $P$ : security policy  
set of executions  $\rightarrow$  Boolean

$S$  is safe iff  $P(\Sigma_S)$  is true.

- Execution Traces



- Security Policies



## ○ Security Policies: Examples

- **Access Control policies** specify that no execution may operate on certain resources such as files or sockets, or invoke certain system operations.
- **Availability policies** specify that if a program acquires a resource during an execution, then it must release that resource at some (arbitrary) later point in the execution.
- **Bounded Availability policies** specify that if a program acquires a resource during an execution, then it must release that resource by some fixed point later in the execution

