

# LANGUAGE BASED SECURITY (LBT)

## SECURE COMPILATION

Chiara Bodei, Gian-Luigi Ferrari

Lecture May, 15 2024 [reloaded]

Lecture May, 17 2024



# Outline

Compiler Correctness



Secure compilation

# Outline



Motivating example  
Overview  
Security via program equivalences  
Fully-abstract compilation  
Fully abstract compilation in practice

# Secure compilation

- What does it mean for a compiler to be secure?

# Example

Rust

$P_1$

$P_2$

...

$P_n$

---

Assembler

$P$

$\llbracket P_1 \rrbracket$

$\llbracket P_2 \rrbracket$

...

$\llbracket P_n \rrbracket$

$P'$

# Example

Rust

`y = &mut`

$P_1$

$P_2$

...

$P_n$

---

Assembler

$P$

$[[P_1]]$

$[[P_2]]$

...


$[[P_n]]$

$P'$


# Example

used linearly

Rust




Assembler




# Linearity

used linearly

Rust



Assembler







# Linearity

used linearly

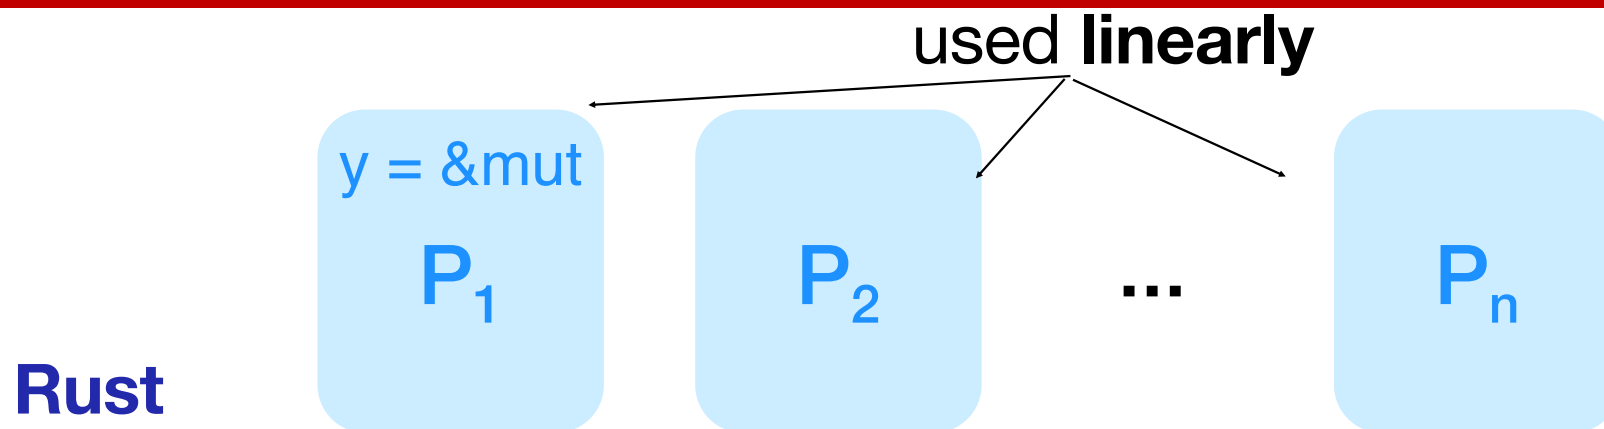
Rust




Assembler



# Possible violations of linearity




## Assembler




# Possible violations of linearity

Preserve the security properties of




Assembler




# Preservation of linearity

Preserve the security properties of

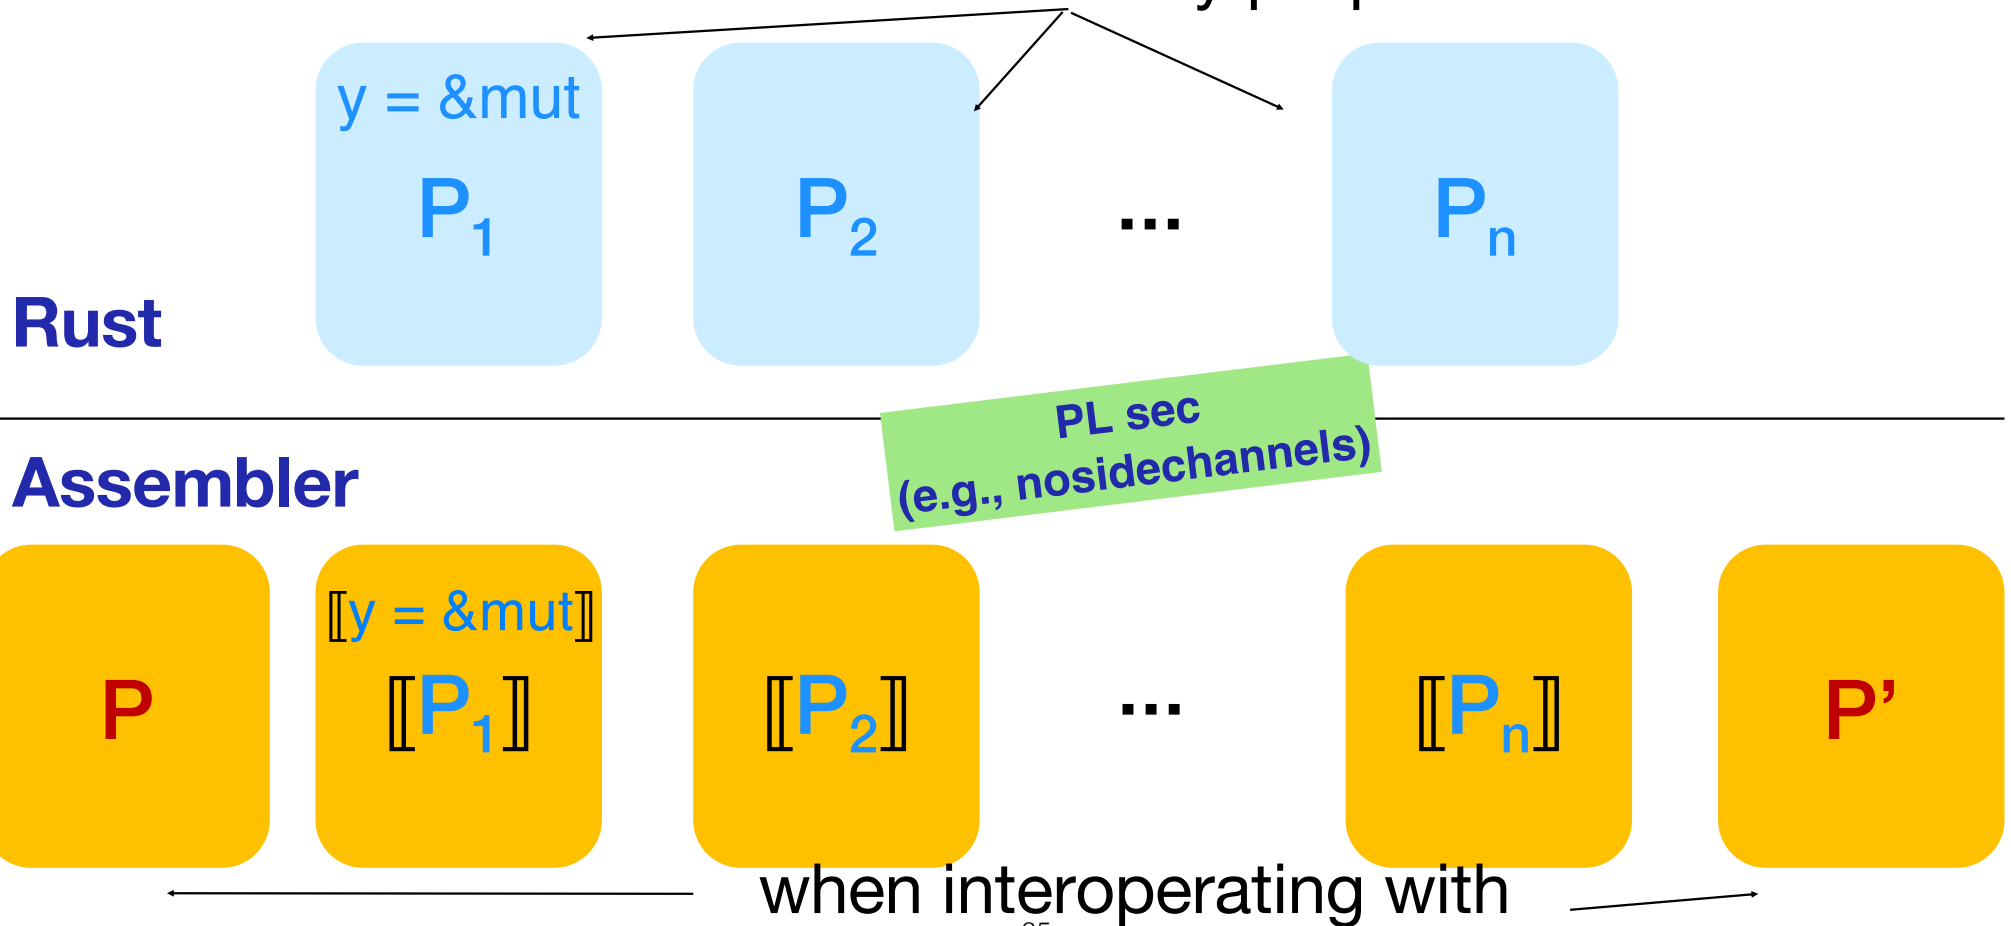


**Assembler**



# Preservation of linearity

Preserve the security properties of




# Threat model hint

What we are  
interested  
in protecting

**Rust**

Preserve the security properties of




**Assembler**




when interoperating with

# Threat model hint




## Assembler



# Correct compilation

Rust




Assembler





# Correct compilation

Rust




Assembler




← respect linearity →

# Secure compilation

Enable source-level security reasoning



Assembler



# Secure compilation

- What does it mean for a compiler  $\llbracket \cdot \rrbracket_T^S$  to be secure?
- Does a given compiler\* preserve the security properties of the source programs?
- Is important this issue?
- What does it mean to preserve security properties? We focus on **formal** answers
- Intuitively, what is secure in the source must be **as** secure in the target

# Correctness vs security

```
int n = some_pt->n;  
if (some_pt == NULL)  
    // Some code  
use (n)
```



```
int n = some_pt->n;  
use (n)
```

```
pin := read_secret();  
if (check(pin))  
    // OK!  
  
pin := 0; // overwrite the pin
```



```
pin := read_secret();  
if (check(pin))  
    // OK!
```

# Abstraction issues

- A high-level language provides a variety of **abstractions** and **mechanisms** (e.g., types, modules, automatic memory management) that enforce good programming
- Unfortunately, most target languages **cannot preserve** the abstractions of their source-level counterparts

# Abstraction issues (cont.)


- The source-level **abstractions** can be used to enforce security properties
- when compiled (and linked with adversarial target code) these abstractions are NOT enforced
- Unfortunately, **discrepancy** between what **abstractions** the source language offers and what abstraction the **target language** has, make target language vulnerable to attacks

# Abstractions\*

- Programming **abstractions** are not preserved by compilers (linkers, etc) (security is an abstraction)
- What does **preserving** abstractions mean?
- **Secure compilation** is an emerging research field concerned with the design and the implementation of compilers that preserve source-level security properties at the object level.

\*Marco Patrignani, Amal Ahmed, Dave Clarke. *Formal Approaches to Secure Compilation. A Survey of Fully Abstract Compilation and Related Work*. ACM Computing surveys, 2019.

# Outline



Motivating example  
Overview  
Security via program equivalences  
Fully-abstract compilation  
Fully abstract compilation in practice



# Example: compilation

```
1 package Bank;
2
3 public class Account{
4     private int balance = 0;
5
6     public void deposit( int amount ) {
7         this.balance += amount;
8     }
9 }
```

Listing 1. Example of Java source code.

Vulnerability addressed  
at the beginning of our course

# Example: compilation

```
1 package Bank;
2
3 public class Account{
4     private int balance = 0;
5
6     public void deposit( int amount ) {
7         this.balance += amount;
8     }
9 }
```

No access to balance  
from outside

Listing 1. Example of Java source code.

# Example: compilation

```
1 package Bank;
2
3 public class Account{
4     private int balance = 0;
5
6     public void deposit( int amount ) {
7         this.balance += amount;
8     }
9 }
```

No access to balance  
from outside  
Enforced by the language

Listing 1. Example of Java source code.

# Example: compilation

```
1 package Bank;
2
3 public class Account{
4     private int balance = 0;
5
6     public void deposit( int amount ) {
7         this.balance += amount;
8     }
9 }
```

No access to balance  
from outside  
Enforced by the language

Listing 1. Example of Java source code.

```
1 typedef struct account_t {
2     int balance = 0;
3     void ( *deposit ) ( struct Account*, int ) = deposit_f;
4 } Account;
5
6 void deposit_f( Account* a, int amount ) {
7     a->balance += amount;
8     return;
9 }
```

Listing 2. C code obtained from compiling the Java code of Listing 1.

# Example: compilation

```
1 package Bank;
2
3 public class Account{
4     private int balance = 0;
5
6     public void deposit( int amount ) {
7         this.balance += amount;
8     }
9 }
```

No access to balance  
from outside  
Enforced by the language

Listing 1. Example of Java source code.

```
1 typedef struct account_t {
2     int balance = 0;
3     void ( *deposit ) ( struct Account*, int ) = deposit_f;
4 } Account;
5
6 void deposit_f( Account* a, int amount )
7     a->balance += amount;
8     return;
9 }
```

Pointer arithmetic in C can lead to  
security violation: undesired access to  
balance

Listing 2. C code obtained from compiling the Java code of Listing 1.

# Example: compilation

- When the Java code interacts with other Java code, the latter cannot access the contents of `balance`, since it is a private field
- However, when the Java code is compiled into the C code and then interacts with arbitrary C code, the latter can access the contents of `balance` by doing simple pointer arithmetic
- Given a pointer to a C `Account struct`, an attacker can add the size (in words) of an `int` to it and read the contents of `balance`, effectively violating a **confidentiality** property that the source program had

# Why?

This **violation** occurs because the **source-level abstractions** used to enforce security properties **is not preserved by the target languages**

# Source-Level Abstractions and Target-Level Attacks

- There are **several examples of source-level security properties** that can be **violated** by target-level attackers that show the security relevance of compilation
- The **capabilities of an attacker vary** depending on the target language considered (typed/untyped,...)
- We will present some examples of the **relevant threats** that a secure compiler needs to mitigate



# Threats: confidentiality

```
1 private secret : Int = 0;  
2  
3 public setSecret() : Int {  
4     secret = 1;  
5     return 0;  
6 }
```

# Threats: confidentiality

```
1 private secret : Int = 0;  
2  
3 public setSecret() : Int {  
4   secret = 1;  
5   return 0;  
6 }
```

No access to balance from outside

BUT if in the target language  
locations are identified by nat numbers  
then the address of secret can be read,  
by dereferencing the number

# Threats: integrity

```
1 public proxy( callback : Unit → Unit ) : Int {  
2   var secret = 1;  
3   callback();  
4   return 0;  
5 }
```

# Threats: integrity

```
1 public proxy( callback : Unit → Unit ) : Int {  
2   var secret = 1;  
3   callback();  
4   return 0;  
5 }
```

The variable `secret` is inaccessible to the code in the callback function at the source level

If the target language can manipulate the call stack, it can access the `secret` variable and change its value

# Threats: memory size

```
1 public kernel( n : Int, callback : Unit →Unit ) : Int {  
2   for (i = 0 to n){  
3     new Object();  
4   }  
5   callback();  
6   // security-relevant code  
7   return 0;  
8 }
```

# Threats: memory size

```
1 public kernel( n : Int, callback : Unit → Unit ) : Int {  
2   for (i = 0 to n){  
3     new Object();  
4   }  
5   callback();  
6   // security-relevant code  
7   return 0;  
8 }
```

If the target language can allocate only  $n$  objects and callback allocates another object, then the security relevant code will not be executed

# Threats: deterministic memory allocation

```
1 public newObjects( ) : Object {  
2     var x = new Object();  
3     var y = new Object();  
4     return x;  
5 }
```

# Threats: deterministic memory allocation

```
1 public newObjects( ) : Object {  
2   var x = new Object();  
3   var y = new Object();  
4   return x;  
5 }
```

At source level, object `y` is inaccessible.

A target level attacker that knows the memory allocation order and can guess where an object will be allocated and influence its memory contents



# Threats: well-typedness

```
1 class Pair {  
2   private first, second : Obj = null;  
3   public getFirst(): Obj {  
4     return this.first;  
5   }  
6 }  
7 class Secret {  
8   private secret : Int = 0;  
9 }  
10 object o : Secret
```

# Threats: well-typedness

```
1 class Pair {  
2   private first, second : Obj = null;  
3   public getFirst(): Obj {  
4     return this.first;  
5   }  
6 }  
7 class Secret {  
8   private secret : Int = 0;  
9 }  
10 object o : Secret
```

- The [Pair class](#) provides a way to store a pair of objects, but it is incomplete as there are no methods to set its values. It only has a method to retrieve the value of first
- The [Secret class](#) does not expose any methods to access or modify its secret value
- The [o object](#) of type Secret is declared but not utilized

# Threats: well-typedness

```
1 class Pair {  
2   private first, second : Obj = null;  
3   public getFirst(): Obj {  
4     return this.first;  
5   }  
6 }  
7 class Secret {  
8   private secret : Int = 0;  
9 }  
10 object o : Secret
```

At target level, an attacker can call `getFirst()` with current object `o`; this will return the secret field, since fields are accessed by offset in untyped assembly

- An attacker, aware of the memory layout of the Secret class, could determine the offset of the secret field within an instance of the Secret class in memory.
- By performing a method call on the `o` object of type Secret, the attacker could potentially access the memory location corresponding to the secret field using the offset obtained earlier.
- This would effectively bypass the access control mechanisms provided by Java and allow the attacker to read the value of the secret field, despite it being declared as private.

# Threats: information flow

```
1 public isZero( value : Inth ) : Intl {  
2   if ( value = 0 ) {  
3     return 1  
4   }  
5   return 0  
6 }
```

Listing 4. Example code with indirect information flow.

# Threats: information flow

```
1 public isZero( value : Inth ) : Intl {  
2   if ( value == 0 ) {  
3     return 1  
4   }  
5   return 0  
6 }
```

The attacker can detect whether value is 0 or not by observing the output. The target language that doesn't prevent information flow cannot withstand these leaks

# Secure compilation

- **[Formally] secure compilation** studies compilers that preserve the security properties of source languages in their compiled, target level counterparts
- What does it mean to **preserve security properties** across compilation?
- Roughly, **having that something secure in the source is still secure in the target**

# Outline



Motivating example  
Overview  
Security via program equivalences  
Fully-abstract compilation  
Fully abstract compilation in practice

# Program equivalences

- A possible way to know what is secure in a program is by exploiting **program equivalences**



# Are these programs equivalent?

<pre>1 public Bool getTrue( x : Bool ) 2   return true;</pre>	. P1	
<pre>1 public Bool getTrue( x : Bool ) 2   return x or true;</pre>	. P2	
<pre>1 public Bool getTrue( x : Bool ) 2   return x and false;</pre>	. P3	
<pre>1 public Bool getTrue( x : Bool ) 2   return false;</pre>	. P4	
<pre>1 public Bool getFalse( x : Bool ) 2   return x and true;</pre>	. P5	

# Are these programs equivalent?

```
1 public Bool getTrue( x : Bool )  
2   return true;
```

```
1 public Bool getTrue( x : Bool )  
2   return x or true;
```

```
1 public Bool getTrue( x : Bool )  
2   return x and false;
```

```
1 public Bool getTrue( x : Bool )  
2   return false;
```

```
1 public Bool getFalse( x : Bool )  
2   return x and true;
```

• P1  
• P2

) =

• P3  
• P4

) =

• P5

# Program equivalences

A possible way to know what is secure in a program is by exploiting **program equivalences**

Roughly, two programs are **equivalent** when they **behave the same** even if they are different (same semantics and possibly different syntax)... **in a way that respects the security property**

- **contextual equivalence**
- **observational equivalence** (e.g., for non-interference property)
- **timing/resource-sensitive equivalence** (e.g., security of constant-time code)

# Confidentiality as equivalence

```
1 private secret : Int = 0;  
2  
3 public setSecret( ) : Int {  
4   secret = 0;  
5   return 0;  
6 }
```

If the two snippets are equivalent  
then secret is confidential

```
1 private secret : Int = 0;  
2  
3 public setSecret( ) : Int {  
4   secret = 1;  
5   return 0;  
6 }
```

# Confidentiality as equivalence

```
1 private secret : Int = 0;  
2  
3 public setSecret( ) : Int {  
4   secret = 0;  
5   return 0;  
6 }
```

If the two snippets are equivalent  
then secret is confidential

```
1 private secret : Int = 0;  
2  
3 public setSecret( ) : Int {  
4   secret = 1;  
5   return 0;  
6 }
```

With a Java-like semantics, secret is  
never accessed from outside  
With a C-like semantics, secret can be  
accessed from outside

# Integrity as equivalence

```
1 public proxy( callback : Unit → Unit )  
  : Int {  
2   var secret = 0;  
3   callback();  
4   if ( secret == 0 ) {  
5     return 0;  
6   }  
7   return 1;  
8 }
```

If the two snippets are equivalent  
then secret cannot be modified  
during the callback

```
1 public proxy( callback : Unit → Unit )  
  : Int {  
2   var secret = 0;  
3   callback();  
4  
5   return 0;  
6  
7  
8 }
```

# Integrity as equivalence

```
1 public proxy( callback : Unit → Unit )  
  : Int {  
2   var secret = 0;  
3   callback();  
4   if ( secret == 0 ) {  
5     return 0;  
6   }  
7   return 1;  
8 }
```

If the two snippets are equivalent  
then secret cannot be modified  
during the callback

```
1 public proxy( callback : Unit → Unit )  
  : Int {  
2   var secret = 0;  
3   callback();  
4  
5   return 0;  
6  
7  
8 }
```

When callback() is invoked,  
secret is on the stack

# Unbounded memory size as equivalence

```
1 public kernel( n : Int, callback : Unit  
    → Unit ) : Int {  
2   for (Int i = 0; i < n; i++){  
3     new Object();  
4   }  
5   callback();  
6   // security-relevant code  
7   return 0;  
8 }
```

```
1 public kernel( n : Int, callback : Unit  
    → Unit ) : Int {  
2  
3  
4  
5   callback();  
6   // security-relevant code  
7   return 0;  
8 }
```



# Unbounded memory size as equivalence

```
1 public kernel( n : Int, callback : Unit  
    → Unit ) : Int {  
2   for (Int i = 0; i < n; i++){  
3     new Object();  
4   }  
5   callback();  
6   // security-relevant code  
7   return 0;  
8 }
```

If the target language can  
allocate only  $n$  objects  
and callback allocates another object,  
then the security  
relevant code will not be executed

```
1 public kernel( n : Int, callback : Unit  
    → Unit ) : Int {  
2  
3  
4  
5   callback();  
6   // security-relevant code  
7   return 0;  
8 }
```

# Memory allocation as equivalence

```
1 public newObjects( ) : Object {  
2     var x = new Object();  
3     var y = new Object();  
4     return x;  
5 }
```

```
1 public newObjects( ) : Object {  
2     var x = new Object();  
3     var y = new Object();  
4     return y;  
5 }
```

# Memory allocation as equivalence

```
1 public newObjects( ) : Object {  
2   var x = new Object();  
3   var y = new Object();  
4   return x;  
5 }
```

```
1 public newObjects( ) : Object {  
2   var x = new Object();  
3   var y = new Object();  
4   return y;  
5 }
```

If the two snippets are equivalent  
then the memory order is  
invisible to an attacker

# Question

Does compilation transform equivalent source-level components into equivalent target-level ones?

# Hence

- The assumption is that program equivalence capture security properties of source code

# Question

How to express program equivalence?

# Which program equivalence?

How to express program equivalence?

## Contextual equivalence

Two programs are **contextually equivalent** if no matter what external observer interacts with them that observer cannot distinguish the programs

# Contextual equivalence

How to express program equivalence?

## Contextual equivalence

Two programs are equivalent if no matter what context/external observer interacts with them that observer cannot distinguish the programs

$$P_1 \simeq_{ctx} P_2 = \forall \mathcal{C}. \mathcal{C}(P_1) \downarrow \Leftrightarrow \mathcal{C}(P_2) \downarrow$$

↓ refers to termination



# Contextual equivalence

The **external observer**  $C$  is generally called **context**

- it is a program, written in the same language as  $P_1$  and  $P_2$
- it is the same program  $C$  interacting with both  $P_1$  and  $P_2$  in two different runs
- so, it cannot express out of language attacks (e.g., side channels)
- interaction means link and run together (like a library)

# Contextual equivalence

**Distinguishing** means: terminate with different values

- the observer basically asks the question: is this program  $P_1$ ?
- if the observer can find a way to distinguish  $P_1$  from  $P_2$ , it will return true, otherwise false
- often, we use divergence and termination as opposed to this boolean termination

# Contexts: an example

**Partial programs:** sequences of assignments of expressions to locations.

- **Expressions:** combination of arithmetic operators and variables  $a_0, a_1, \dots$
- **Locations:**  $X_0, X_1, \dots$
- **Contexts:** non-empty lists of natural numbers
- Linking a context  $C$  and a partial program  $P$  gives a whole program  $C[P]$  in which the variables in expressions of  $P$  are initialized with the information provided by  $C$

**P**

$X_0 := (a_0 \times a_1) \times a_2$

$X_1 := a_0$

If  $C = [2, 3, 7]$ , then  $C[P]$  is

$X_0 := (2 \times 3) \times 7$

$X_1 := 2$

# Context

- A **context**  $C$  is a program with a **hole** (denoted by  $[.]$ ), which can be filled by a component  $P$ , generating the whole program  $C[P]$  to be executed
- Plugging a component in a **context** makes the program whole, so its behaviour can be observed via its operational semantics

# Contextual equivalence

In a simple functional language

$$\lambda x. x * 2 \approx \lambda x. x + x$$

# Contextual equivalence

In a simple functional language

$$\lambda x. x * 2 \approx \lambda x. x + x$$

in Javascript

`function(x){ return x * 2; }  $\approx$  function(x){ return x + x; }?`

# Contextual equivalence

In a simple functional language

$$\lambda x. x * 2 \approx \lambda x. x + x$$

In Javascript

$$\text{function}(x)\{\text{return } x * 2; \} \approx \text{function}(x)\{\text{return } x + x; \}?$$

Is there a program context that can distinguish them?

# Contextual equivalence

In a simple functional language

$$\lambda x. x * 2 \approx \lambda x. x + x$$

In Javascript

`function(x){ return x * 2; }  $\approx$  function(x){ return x + x; }?`

Is there a program context that can distinguish them?

Yes: `([·]).toString()`

Troublesome for optimizations

`toString()` returns the source code of the function



# Confidentiality as equivalence

```
1 private secret : Int = 0;  
2  
3 public setSecret( ) : Int {  
4   secret = 0;  
5   return 0;  
6 }
```

```
1 private secret : Int = 0;  
2  
3 public setSecret( ) : Int {  
4   secret = 1;  
5   return 0;  
6 }
```

```
// Observer P in Java  
2 public static isItP1( ) : Bool  
{  
3   Secret.getSecret();  
4   ...  
5 }
```

P cannot tell the difference!

# Confidentiality as equivalence

```
1 private secret : Int = 0;  
2  
3 public setSecret( ) : Int {  
4   secret = 0;  
5   return 0;  
6 }
```

```
1 private secret : Int = 0;  
2  
3 public setSecret( ) : Int {  
4   secret = 1;  
5   return 0;  
6 }
```

```
1 // Observer P in Java  
2 public static isItP1( ) : Bool  
3 {  
4   Secret.getSecret();  
5   ...  
6 }
```

P cannot tell the difference!  
There is no getSecret() method

# Confidentiality as equivalence

```
1 typedef struct secret { // P1
2   int secret = 0;
3   void (*setSec) (struct Secret*) = setSec;
4 } Secret;
5 void setSec(Secret* s) {s→secret = 0; return;}
```

```
1 typedef struct secret { // P2
2   int secret = 0;
3   void (*setSec) (struct Secret*) = setSec;
4 } Secret;
5 void setSec(Secret* s) {s→secret = 1; return;}
```

# Confidentiality as equivalence

```
1 typedef struct secret { // P1
2   int secret = 0;
3   void (*setSec) (struct Secret*) = setSec;
4 } Secret;
5 void setSec(Secret* s) {s->
```

```
1 typedef struct secret { //
2   int secret = 0;
3   void (*setSec) (struct Sec
4 } Secret;
5 void setSec(Secret* s) {s->
```

1 // Observer P in C

2 int isItP1(){

3 struct Secret x;

4 sec = &x + sizeof(int);

5 if \*sec == 0 then return true else return false

6 }

P can see the difference!

The two programs are inequivalent

# Security violations

## Inequivalences as security violations

if the target programs are not equivalent then the intended security property is violated

# Security preservation

What does it mean to preserve security properties across compilation?

# Security preservation

What does it mean to preserve security properties across compilation?

Given source equivalent programs (which have a security property), compile them into equivalent target programs

# Security preservation

What does it mean to preserve security properties across compilation?

Given source equivalent programs (which have a security property), compile them into equivalent target programs, provided that:

- the security property is captured in the source by program equivalence

Being equivalent in the target means contextual equivalence w.r.t. **target observers** (i.e., target programs), i.e., the **attackers** in this setting



# Secure compilation

Recall that

- **Attackers** are modelled as the **environment programs** to be checked interact with (link and run together)
- **Attackers can act** and not only observe the program behaviour

# Partial programs

Note that

- For **correct compilation**, we considered **whole programs**
- **Secure compilation** is instead concerned with the security of **partial programs** (or components) that are linked together with an environment (or context)

# Outline



Motivating example  
Overview  
Security via program equivalences  
Fully-abstract compilation  
Fully-abstract compilation in practice

# Formal (secure compilation): full abstraction

A compiler  $\llbracket \cdot \rrbracket$  is **fully abstract** when it translates equivalent source-level components into equivalent target-level ones

$\forall P_1, P_2$

$$P_1 \simeq_{\text{ctx}} P_2 \Rightarrow \llbracket P_1 \rrbracket \simeq_{\text{ctx}} \llbracket P_2 \rrbracket$$

If the target programs are not equivalent, then the intended security property is violated

Is it enough?

# Formal (secure compilation): full abstraction

A compiler  $\llbracket \cdot \rrbracket$  is **fully abstract** when it translates equivalent source-level components into equivalent target-level ones

$\forall P_1, P_2$

$$P_1 \simeq_{\text{ctx}} P_2 \Rightarrow \llbracket P_1 \rrbracket \simeq_{\text{ctx}} \llbracket P_2 \rrbracket$$

If the target programs are not equivalent, then the intended security property is violated

Is it enough? No, it is not

An empty translation would fit

# Correctness is needed

We need the compiler also to be **correct**, i.e.,

$\forall P_1, P_2$

$$P_1 \simeq_{\text{ctx}} P_2 \Leftarrow \llbracket P_1 \rrbracket \simeq_{\text{ctx}} \llbracket P_2 \rrbracket$$

# Fully Abstract compilation

Then we have

$\forall P_1, P_2$

$$P_1 \simeq_{\text{ctx}} P_2 \Leftrightarrow \llbracket P_1 \rrbracket \simeq_{\text{ctx}} \llbracket P_2 \rrbracket$$

## Correctness

**Reflection of  $\simeq$**

$$P_1 \simeq P_2 \Leftarrow \llbracket P_1 \rrbracket \simeq \llbracket P_2 \rrbracket$$

The compiler outputs behave as their source-level counterparts

## Security

**Preservation of  $\simeq$**

$$P_1 \simeq P_2 \Rightarrow \llbracket P_1 \rrbracket \simeq \llbracket P_2 \rrbracket$$

The source-level abstractions are not violated in the target-level output

# Fully Abstract compilation

Then we have

$\forall P_1, P_2$

$$P_1 \simeq_{\text{ctx}} P_2 \Leftrightarrow \llbracket P_1 \rrbracket \simeq_{\text{ctx}} \llbracket P_2 \rrbracket$$

$\forall P_1, P_2.$

$$\forall \mathcal{C}.. \mathcal{C}(P_1) \simeq_{\text{ctx}} \mathcal{C}(P_2) \Rightarrow \forall \mathcal{C}.. \mathcal{C}(\llbracket P_1 \rrbracket) \simeq_{\text{ctx}} \mathcal{C}(\llbracket P_2 \rrbracket)$$



# Correctness is needed

We need the compiler to be **correct**

$$\forall P_1, P_2. P_1 \simeq_{\text{ctx}} P_2 \Leftarrow \llbracket P_1 \rrbracket \simeq_{\text{ctx}} \llbracket P_2 \rrbracket$$

- The **contrapositive**\* is actually easier to understand:  
if the source components were not equivalent then the target components would have to be different
- If this did not hold, then the compiler could take different source components and compile them to the same target component
- $P \Rightarrow Q$  corresponds to  $\neg Q \Rightarrow \neg P$  (If it is raining, then I open my umbrella" — "If I do not open my umbrella, then it is not raining.")

# Correctness is needed

Note that **equivalence** is related to **compiler correctness**, but the two notions do not coincide

$$\forall P_1, P_2. P_1 \simeq_{ctx} P_2 \Leftarrow \llbracket P_1 \rrbracket_{ctx} \llbracket P_2 \rrbracket$$

As long as your compiler produced different target programs for different source programs, all would be fine: e.g., hash the source program and produce target programs that just printed the hash

Different target programs not only for source programs that were observationally different, but even syntactically different!

A pretty bad compiler, and certainly not correct, but it would be equivalence reflecting

# Equivalence preservation

Equivalence preservation is the hallmark of fully abstract compilers

$$\forall P_1, P_2. P_1 \simeq_{ctx} P_2 \Rightarrow \llbracket P_1 \rrbracket \simeq_{ctx} \llbracket P_2 \rrbracket$$

Observers in the target cannot make observations that are not possible to distinguish in the source

If a programming language includes features for information hiding, such as private fields, and two source components are identical except for having different values stored in those private fields, problems can arise if the compiler does not properly preserve the privacy of those fields

# Compiler full abstraction

If two programs are **equivalent** in the source language (i.e., no source context can distinguish them), the two programs obtained by compiling them are equivalent in the target language (i.e., no target context can distinguish them)

A **fully abstract compilation chain protects source-level abstractions** all the way down, ensuring that linked adversarial target code cannot observe more about the compiled program than what some linked source code could about the source program

A **fully abstract compiler does not eliminate source-level security flaws**, it only introduces no more vulnerabilities at the target-level

# Compiler full abstraction

- **Equivalence-preserving compilation** considers all target-level contexts when establishing indistinguishability, so it captures the power of an attacker operating at the level of the target language
- **Full abstraction allows for source-level reasoning**: the programmer need not be concerned with the behaviour of target-level code (attackers) and can focus only potential source-level behaviors when reasoning about safety and security properties of their code

# Compiler full abstraction

- FA only preserves security property expressed as program equivalence
- FA is not the silver bullet: there are shortcomings of fully abstract compilation
  - Difficult to (dis)-prove a compiler (not) to be FA
  - FA compilers may produce inefficient code
  - Mainstream compilers are not usually FA

# Proving full abstraction

A compiler  $\llbracket \cdot \rrbracket$  is fully abstract if it reflects and preserves the contextual equivalence

$$\text{Preservation} = \forall P_1, P_2 \in S. P_1 \simeq_{\text{ctx}} P_2 \Rightarrow \llbracket P_1 \rrbracket_{\top}^S \simeq_{\text{ctx}} \llbracket P_2 \rrbracket_{\top}^S,$$

$$\text{Reflection} = \forall P_1, P_2 \in S. \llbracket P_1 \rrbracket_{\top}^S \simeq_{\text{ctx}} \llbracket P_2 \rrbracket_{\top}^S \Rightarrow P_1 \simeq_{\text{ctx}} P_2.$$

Recall that  $P_1 \simeq_{\text{ctx}} P_2$  means that  $\forall \mathcal{C}. \mathcal{C}(P_1) \downarrow \Leftrightarrow \mathcal{C}(P_2) \downarrow$

- Both parts are difficult to prove
- Preservation is particularly *tricky* because of  $\simeq$  and  $\forall$  contexts

# Proving preservation

## Preservation

Unfolding context equivalence at the target level:

$\forall P_1, P_2.$

$$P_1 \simeq_{ctx} P_2 \Rightarrow \forall C.. C(\llbracket P_1 \rrbracket) \downarrow \Leftrightarrow C(\llbracket P_2 \rrbracket) \downarrow$$

Contrapositive\*:  $\forall P_1, P_2.$

$$\exists C.. C(\llbracket P_1 \rrbracket) \downarrow \not\Leftrightarrow C(\llbracket P_2 \rrbracket) \downarrow \Rightarrow P_1 \not\leq_{ctx} P_2$$

Unfolding context equivalence at the source level:

$$\exists C.. C(\llbracket P_1 \rrbracket) \downarrow \not\Leftrightarrow C(\llbracket P_2 \rrbracket) \downarrow \Rightarrow \exists C.. C(P_1) \downarrow \not\Leftrightarrow C(P_2) \downarrow$$

\*  $P \Rightarrow Q$  corresponds to  $\neg Q \Rightarrow \neg P$



# Proving preservation

## Preservation

Unfolding context equivalence at the target level:

$\forall P_1, P_2.$

$$P_1 \simeq_{ctx} P_2 \Rightarrow \forall \mathcal{C}.. \mathcal{C}(\llbracket P_1 \rrbracket) \downarrow \Leftrightarrow \mathcal{C}(\llbracket P_2 \rrbracket) \downarrow$$

Contrapositive\*:  $\forall P_1, P_2.$

$$\exists \mathcal{C}.. \mathcal{C}(\llbracket P_1 \rrbracket) \downarrow \not\Leftrightarrow \mathcal{C}(\llbracket P_2 \rrbracket) \downarrow \Rightarrow P_1 \not\leq_{ctx} P_2$$

Preservation amounts to proving that no new vulnerabilities are introduced at the target-level

Unfolding context equivalence at the source level:

$$\exists \mathcal{C}.. \mathcal{C}(\llbracket P_1 \rrbracket) \downarrow \not\Leftrightarrow \mathcal{C}(\llbracket P_2 \rrbracket) \downarrow \Rightarrow \exists \mathcal{C}.. \mathcal{C}(P_1) \downarrow \not\Leftrightarrow \mathcal{C}(P_2) \downarrow$$

\*  $P \Rightarrow Q$  corresponds to  $\neg Q \Rightarrow \neg P$

# Proving preservation

Preservation:

$$\forall P_1, P_2. \forall \mathcal{C}.. \mathcal{C}(P_1) \approx_{\text{ctx}} \mathcal{C}(P_2) \Rightarrow \forall \mathcal{C}.. \mathcal{C}(\llbracket P_1 \rrbracket) \approx_{\text{ctx}} \mathcal{C}(\llbracket P_2 \rrbracket)$$

$$\begin{array}{c} \mathcal{C}(P_1) \approx_{\text{ctx}} \mathcal{C}(P_2) \\ \approx \qquad \qquad \approx \\ \mathcal{C}(\llbracket P_1 \rrbracket) \approx_{\text{ctx}}^? \mathcal{C}(\llbracket P_2 \rrbracket) \end{array}$$

# Security proof schema

$$\forall P_1, P_2. P_1 \simeq_{ctx} P_2 \Rightarrow \llbracket P_1 \rrbracket \simeq_{ctx} \llbracket P_2 \rrbracket$$

Given an arbitrary target-level context  $\mathcal{C}$ , we need to prove that

$$\mathcal{C}(\llbracket P_1 \rrbracket) \simeq_{ctx} \mathcal{C}(\llbracket P_2 \rrbracket)$$

If the source language is strong enough, we can construct a **back-translation**  $\llbracket \cdot \rrbracket$  for any target-level context  $\mathcal{C}$ , obtaining a corresponding valid source-level context  $\mathcal{C}'$  “observational equivalent to”  $\mathcal{C}$

# Back-translations

**Back-translation** of target-level program contexts to behaviourally-equivalent source-level contexts, a sort of compiler in reverse

$[\![\cdot]\!]: S \rightarrow T$

$\llbracket \cdot \rrbracket: T \rightarrow S$

For any target-level context  $\mathcal{C}$  returns a corresponding valid source-level context  $\mathcal{C}$

Constructing such a **back-translation** can be difficult when the source language is not strong enough to embed an encoding of the target language

# Security proof schema

$$\forall P_1, P_2. P_1 \simeq_{ctx} P_2 \Rightarrow \llbracket P_1 \rrbracket \simeq_{ctx} \llbracket P_2 \rrbracket$$

Given an arbitrary target-level context  $\mathcal{C}$ , we need to prove that

$$\mathcal{C}(\llbracket P_1 \rrbracket) \simeq_{ctx} \mathcal{C}(\llbracket P_2 \rrbracket)$$

If the source language is strong enough, we can construct a **back-translation**  $\llbracket \cdot \rrbracket$  for any target-level context  $\mathcal{C}$ , obtaining a corresponding valid source-level context  $\mathcal{C}'$  “observational equivalent to”  $\mathcal{C}$

If  $\mathcal{C}'(P_1) \simeq \mathcal{C}(\llbracket P_1 \rrbracket)$  and  $\mathcal{C}'(P_2) \simeq \mathcal{C}(\llbracket P_2 \rrbracket)$  we can prove that

$$\mathcal{C}(\llbracket P_1 \rrbracket) \simeq_{ctx} \mathcal{C}(\llbracket P_2 \rrbracket)$$

# Proving preservation

Preservation:

$$\forall P_1, P_2. \forall \mathcal{C}.. \mathcal{C}(P_1) \approx_{\text{ctx}} \mathcal{C}(P_2) \Rightarrow \forall \mathcal{C}.. \mathcal{C}(\llbracket P_1 \rrbracket) \approx_{\text{ctx}} \mathcal{C}(\llbracket P_2 \rrbracket)$$

$$\begin{array}{c} \mathcal{C}(P_1) \approx_{\text{ctx}} \mathcal{C}(P_2) \\ \approx \qquad \qquad \approx \\ \mathcal{C}(\llbracket P_1 \rrbracket) \approx_{\text{ctx}}^? \mathcal{C}(\llbracket P_2 \rrbracket) \end{array}$$

# Security proof schema

$$P_1 \simeq_{ctx} P_2$$

$$\simeq$$
$$\simeq$$

$$\llbracket P_1 \rrbracket \stackrel{?}{\simeq}_{ctx} \llbracket P_2 \rrbracket$$

# Security proof schema

$$\begin{array}{ccc}
 P_1 & \simeq_{ctx} & P_2 \\
 \simeq & & \simeq \\
 \llbracket P_1 \rrbracket & \stackrel{?}{\simeq}_{ctx} & \llbracket P_2 \rrbracket \\
 \mathcal{C}(\llbracket P_1 \rrbracket) \downarrow & \stackrel{?}{\Rightarrow} & \mathcal{C}(\llbracket P_2 \rrbracket) \downarrow
 \end{array}$$

Given an arbitrary context  $\mathcal{C}$ , we need to prove that  $\mathcal{C}(\llbracket P_1 \rrbracket) \simeq_{ctx} \mathcal{C}(\llbracket P_2 \rrbracket)$

If  $\llbracket \mathcal{C} \rrbracket = \mathcal{C}'$  such that  $\llbracket \mathcal{C} \rrbracket = \mathcal{C}' \simeq \mathcal{C}$  and  $P_i \simeq \llbracket P_i \rrbracket$  and hence  $\mathcal{C}'(P_1) \simeq \mathcal{C}(\llbracket P_1 \rrbracket)$ , we have done



# Security proof schema

$$P_1 \simeq_{ctx} P_2$$

$$\llbracket \mathbb{C} \rrbracket \simeq \mathbb{C}$$

$$\llbracket \mathbb{C} \rrbracket \simeq \mathbb{C}$$

$$P_1 \simeq \llbracket P_1 \rrbracket$$


$$P_2 \simeq \llbracket P_1 \rrbracket$$

$$\llbracket P_1 \rrbracket \stackrel{?}{\simeq}_{ctx} \llbracket P_2 \rrbracket$$

$$\mathbb{C}(\llbracket P_1 \rrbracket) \downarrow \stackrel{?}{\Rightarrow} \mathbb{C}(\llbracket P_2 \rrbracket) \downarrow$$

Given an arbitrary context  $\mathbb{C}$ , we need to prove that  $\mathbb{C}(\llbracket P_1 \rrbracket) \simeq_{ctx} \mathbb{C}(\llbracket P_2 \rrbracket)$

# Security proof schema



To prove  $\mathcal{C}(\llbracket P_1 \rrbracket) \simeq_{ctx} \mathcal{C}(\llbracket P_2 \rrbracket)$  we exploit the chain of equivalences (1), (2), (3), where (2) comes from the assumed (contextual) equivalence of  $P_1$  and  $P_2$

# Note on equivalences

Some kinds of equivalences may simplify these proofs. E.g., contextual equivalence at the target level can be replaced with trace equivalence if it is proved just as precise

- **Context-based**: relies on the structure of the context
- **Trace-based**: relies on trace semantics

# Note on equivalences

Some kinds of equivalences may simplify these proofs. E.g., contextual equivalence at the target level can be replaced with trace equivalence if it is proved just as precise

- **Context-based**: relies on the structure of the context: when source and target contexts are similar
- **Trace-based**: relies on trace semantics: when there is a large abstraction gap between source and target

# Trace Semantics

We replace context equivalence with something equivalent

- but simpler to reason about
- a semantics that abstracts from the context (observer)
- and still describes the behaviour of a program precisely

A trace semantics

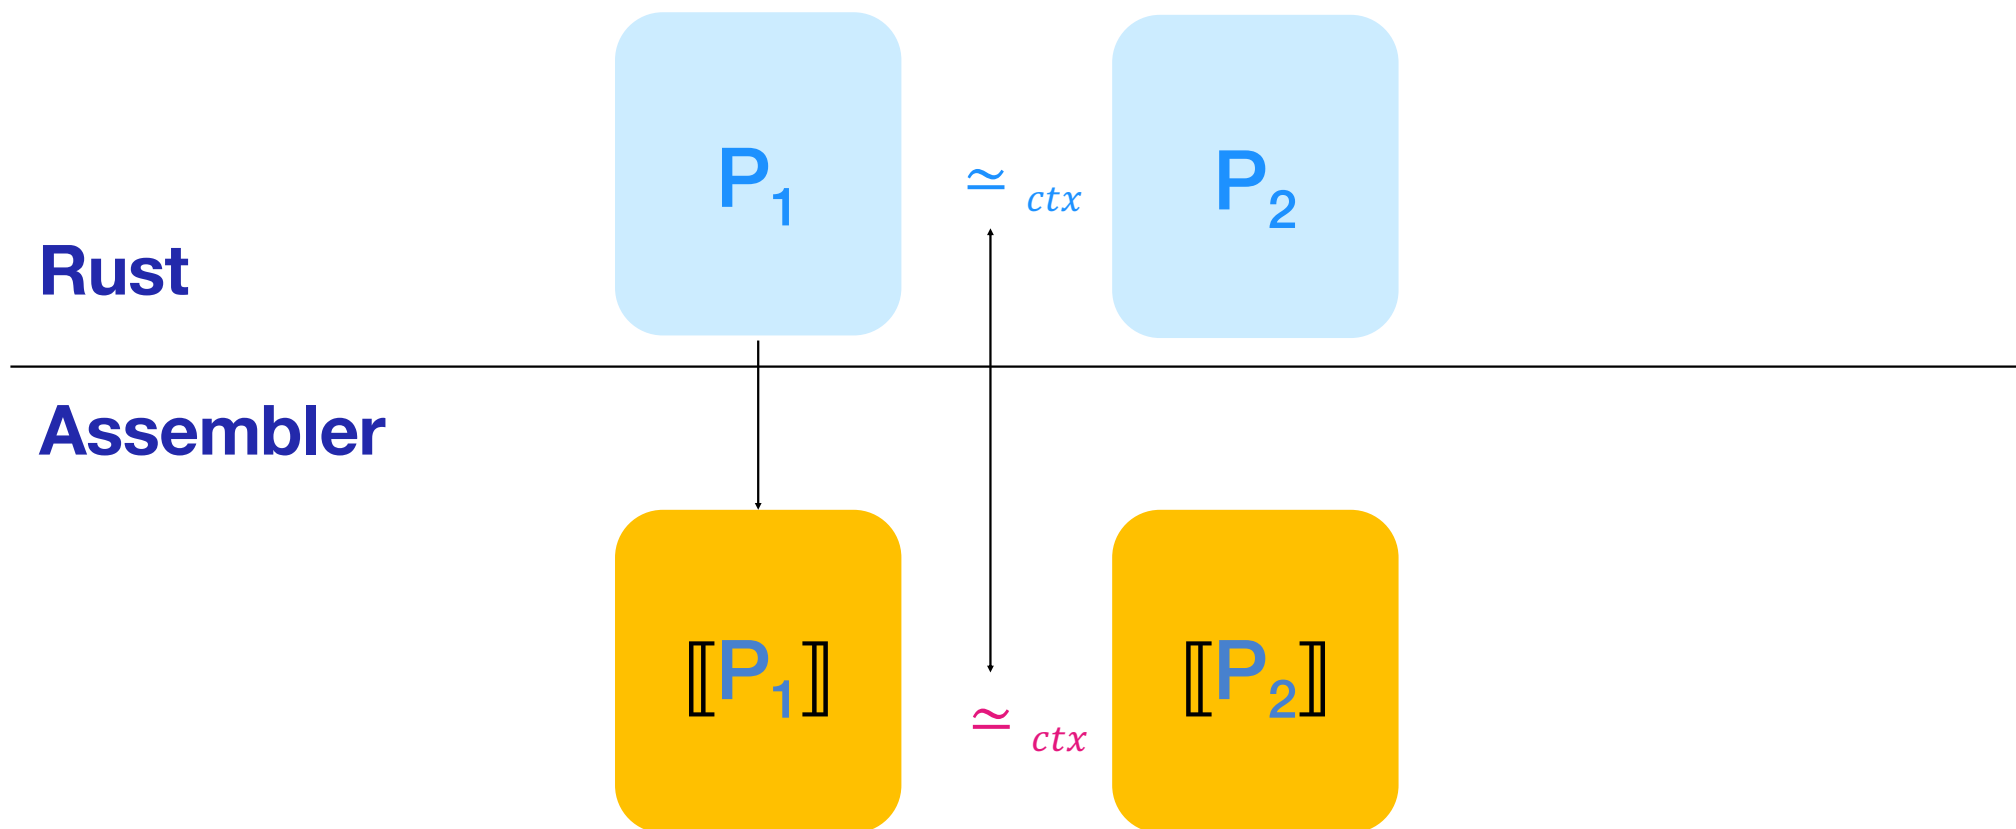
# Trace Semantics

**Trace semantics** for partial programs (component)

- relies on the operational semantics
- denotational: describes the **behaviour** of a component as **sets of traces**
- a **trace** is (typically) a sequence of actions that describe how a component interacts with an observer
- without needing to specify the observer
- Trace semantics come with **trace equivalence**
- Two programs are trace equivalent, when the sets of traces coincide
- We need a trace equivalence that is correct and complete w.r.t. context equivalence

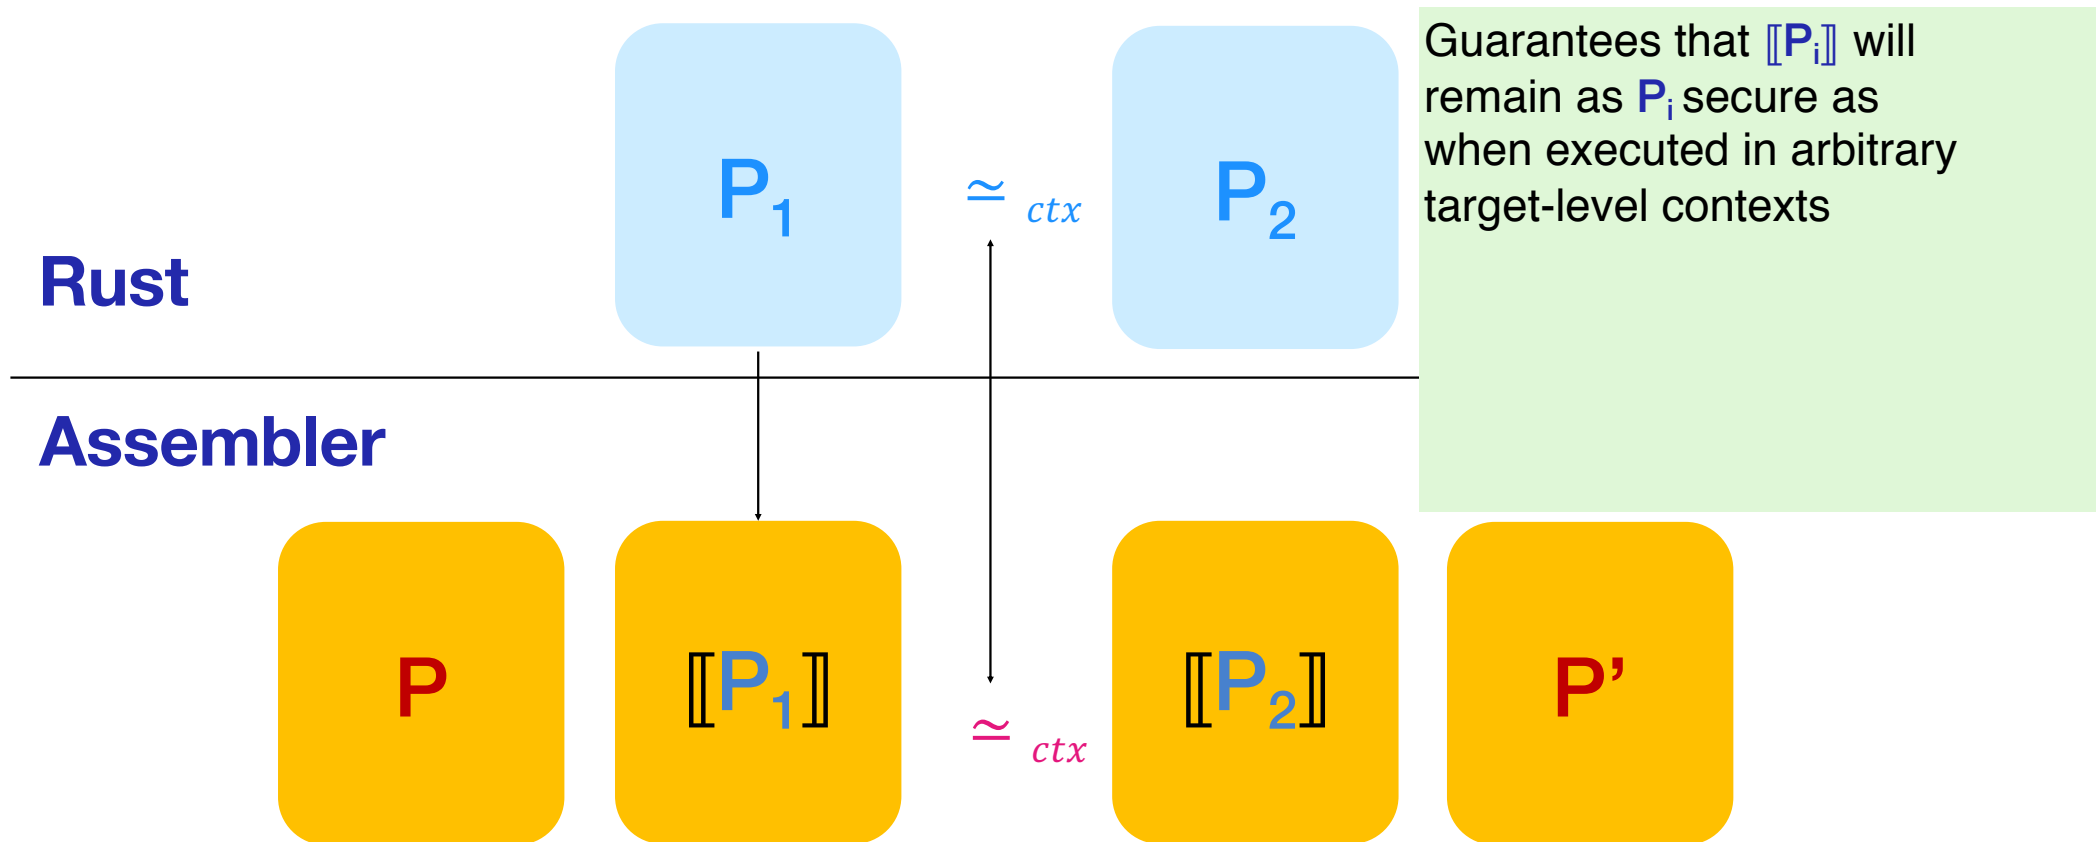
# Full abstract compilation

Preserve and reflect contextual equivalence



# Full abstract compilation

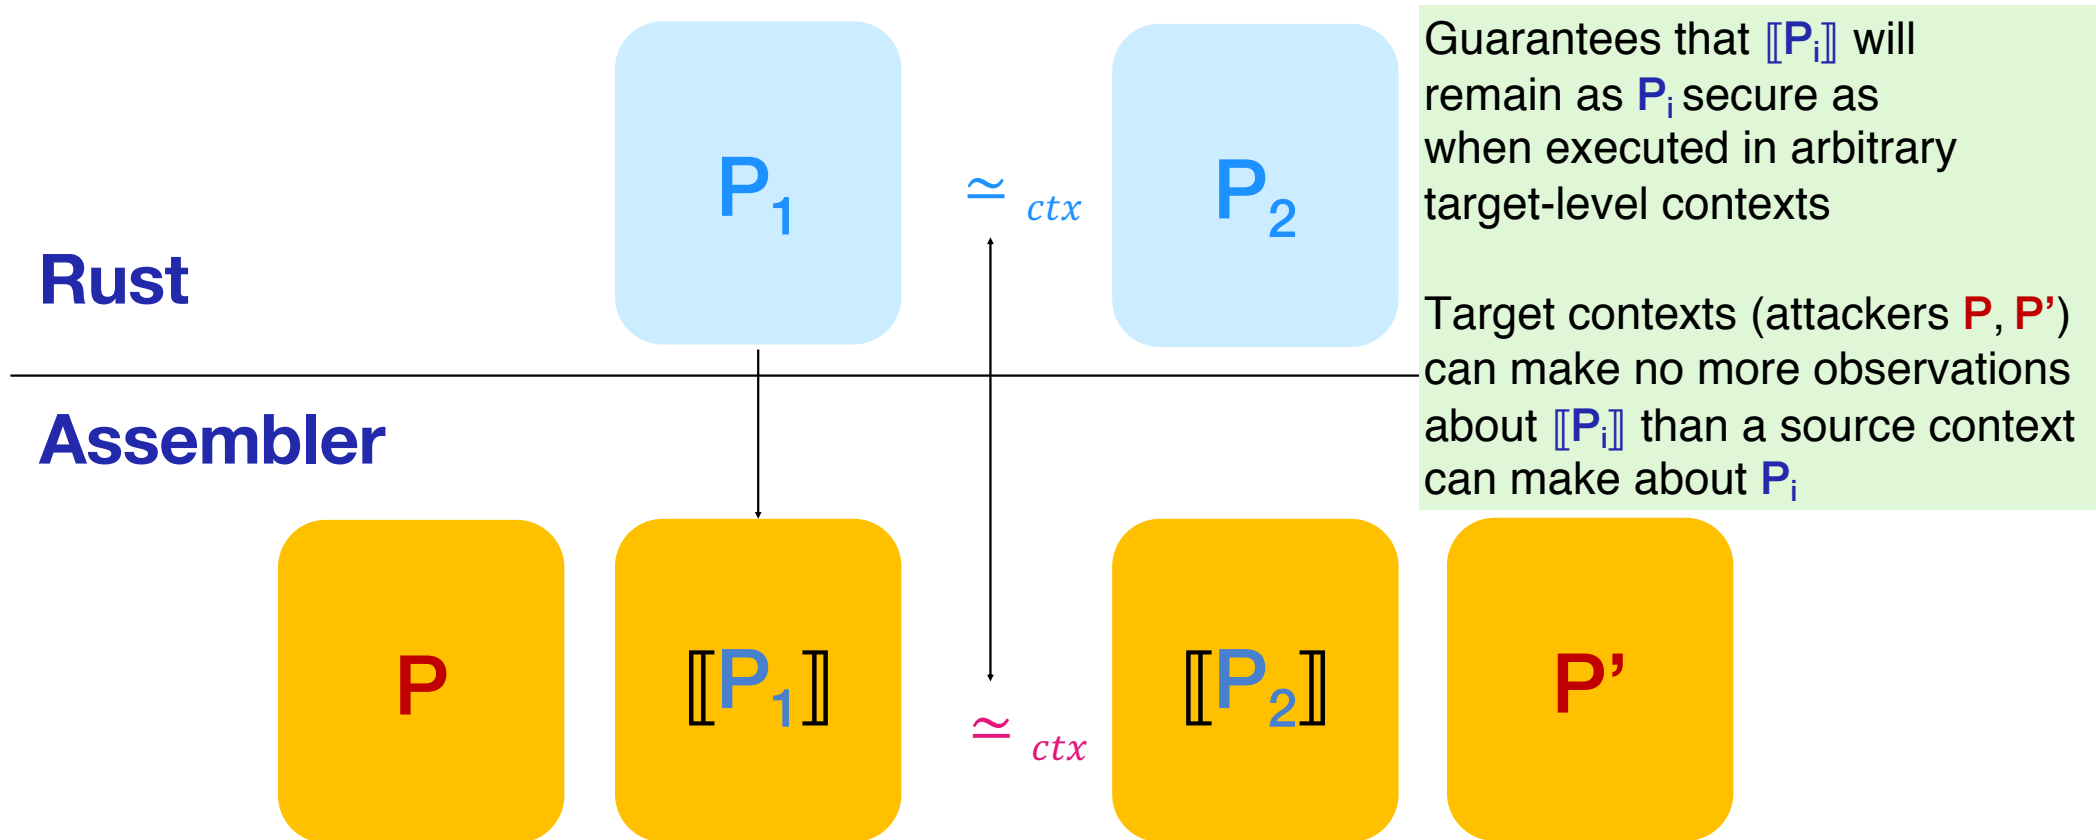
**Preserve** and reflect contextual equivalence





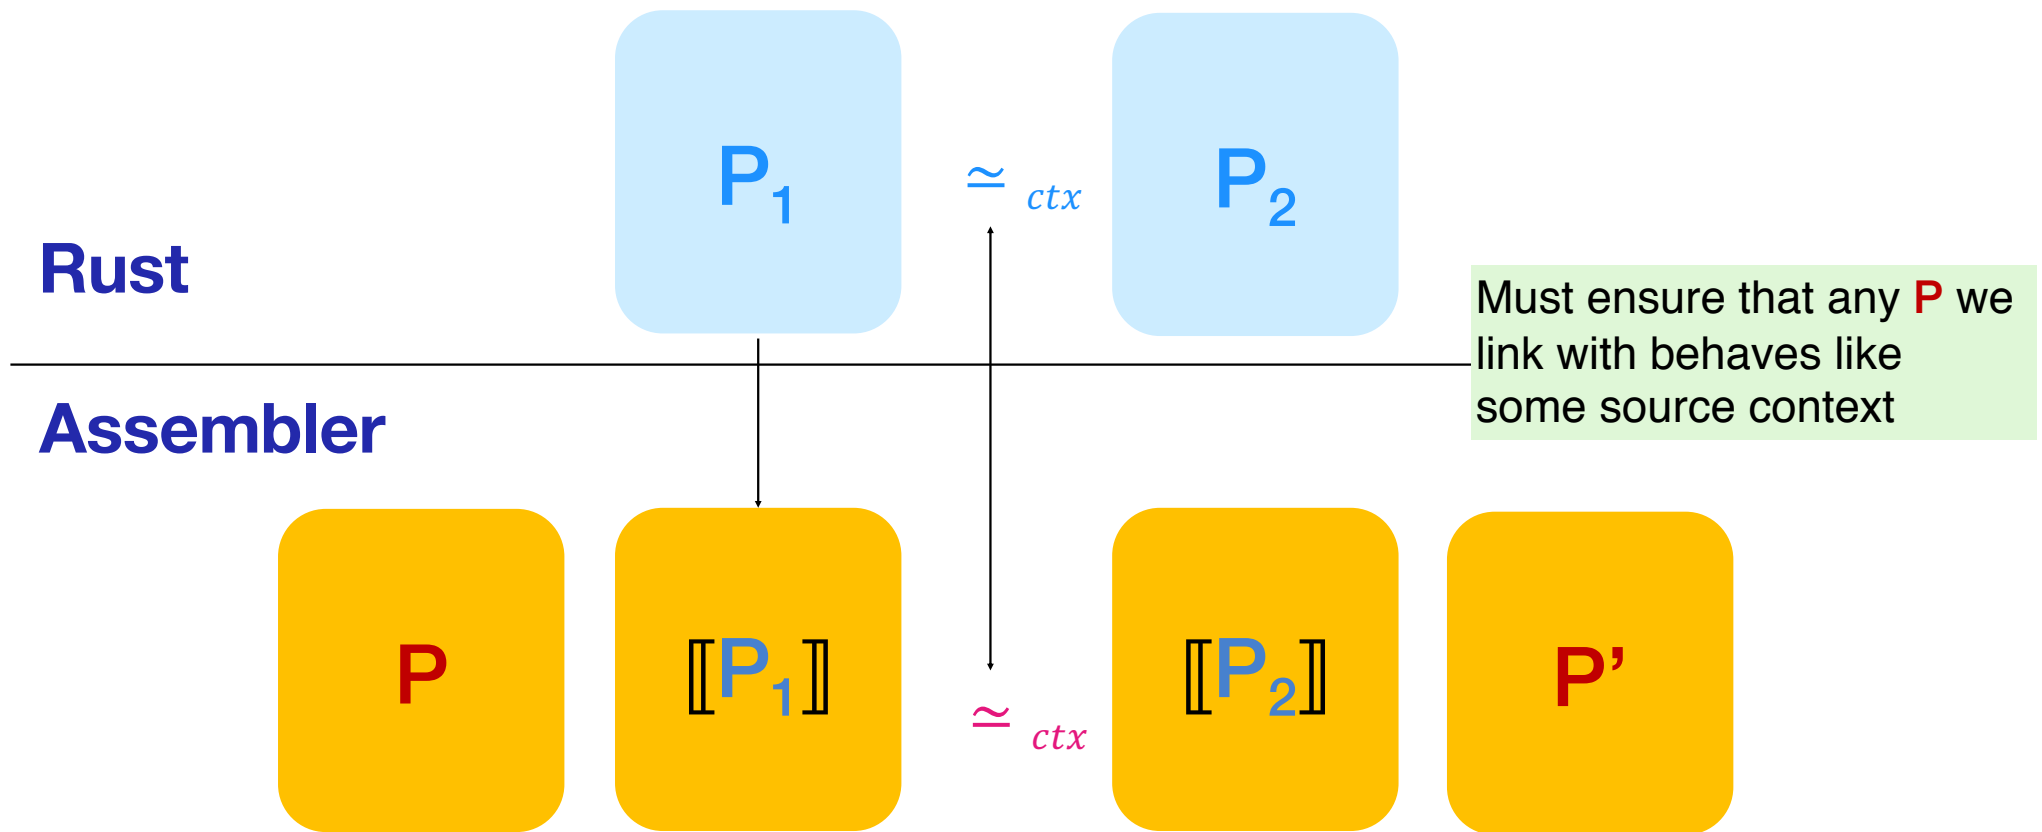
# Full abstract compilation

**Preserve** and reflect contextual equivalence



# Full abstract compilation

**Preserve** and reflect contextual equivalence



# Fully abstract compilation in practice\*

Back to our previous example of compilation from a simple imperative language IMP to a stack-based Virtual Machine VM

We will use **Coq**, a formal proof management system



\*<https://dbp.io/essays/2018-04-19-how-to-prove-a-compiler-fully-abstract.html>

# IMP Arithmetic expressions

Inductive `Expr` : Set := | `Num` :  $\mathbb{Z} \rightarrow \text{Expr}$  | `Plus` :  $\text{Expr} \rightarrow \text{Expr} \rightarrow \text{Expr}$ .

## Evaluation function with 2 constructors

Fixpoint `eval_Expr` (e : Expr) :  $\mathbb{Z}$  := match e with

| `Num` n => n /\*takes n and constructs an expr representing that number \*/

# IMP Arithmetic expressions

Inductive `Expr` : Set := | `Num` :  $\mathbb{Z} \rightarrow \text{Expr}$  | `Plus` :  $\text{Expr} \rightarrow \text{Expr} \rightarrow \text{Expr}$ .

## Evaluation function with 2 constructors

Fixpoint `eval_Expr` (e : Expr) :  $\mathbb{Z}$  := match e with

| `Num` n => n /\*takes n and constructs an expr representing that number \*/

| `Plus` e1 e2 => eval\_Expr e1 + eval\_Expr e2 /\*takes two arguments and constructs an expression representing their sum. \*/

end.

# IMP Arithmetic expressions

Inductive `Expr` : Set := | `Num` :  $\mathbb{Z} \rightarrow \text{Expr}$  | `Plus` :  $\text{Expr} \rightarrow \text{Expr} \rightarrow \text{Expr}$ .

## Evaluation function with 2 constructors

Fixpoint `eval_Expr` (e : Expr) :  $\mathbb{Z}$  := match e with

| `Num` n => n /\*takes n and constructs an expr representing that number \*/

| `Plus` e1 e2 => eval\_Expr e1 + eval\_Expr e2 /\*takes two arguments and constructs an expression representing their sum. \*/

end.

**Example:** `eval_Expr(Plus (Num 1) (Num 2)) = eval_Expr(1) + eval_Expr(2)`  
gives 3 as output

# VM Operations

Inductive Op : Set :=

| Push :  $\mathbb{Z} \rightarrow \text{Op}$  /\* takes an integer value n and constructs an operation to push that value onto the stack \*/

# VM Operations

Inductive  $\text{Op} : \text{Set} :=$

- |  $\text{Push} : \mathbb{Z} \rightarrow \text{Op}$  /\* takes an integer value  $n$  and constructs an operation to push that value onto the stack \*/
- |  $\text{Add} : \text{Op}$  /\* represents addition on the top two elements of the stack \*/



# VM Operations

Inductive  $\text{Op} : \text{Set} :=$

- |  $\text{Push} : \mathbb{Z} \rightarrow \text{Op}$  /\* takes an integer value  $n$  and constructs an operation to push that value onto the stack \*/
- |  $\text{Add} : \text{Op}$  /\* represents addition on the top two elements of the stack \*/
- |  $\text{OpCount} : \text{Op}$ . /\* returns the count of remaining operations on the stack machine and puts that integer on the top of the stack \*/

# VM Operations

Inductive  $Op : Set :=$

|  $Push : \mathbb{Z} \rightarrow Op$  /\* takes an integer value  $n$  and constructs an operation to push that value onto the stack \*/

|  $Add : Op$  /\* represents addition on the top two elements of the stack \*/

|  $OpCount : Op.$  /\* returns the count of remaining operations on the stack machine and puts that integer on the top of the stack \*/

The stack must be empty at the end of execution

# Evaluation of a list of VM operations

Fixpoint `eval_Op` (s : list Z) (ops : list Op) : option Z :=

match (ops, s) with

| ([], [n]) => Some n /\* If the list is empty and there is only one element on the stack, it returns Some n, indicating successful evaluation with the result n \*/

# Evaluation of a list of VM operations

Fixpoint `eval_Op` (s : list Z) (ops : list Op) : option Z :=

match (ops, s) with

| ([], [n]) => Some n /\* If the list is empty and there is only one element on the stack, it returns Some n, indicating successful evaluation with the result n \*/

| (Push z :: rest, \_) => eval\_Op (z :: s) rest /\* If the 1<sup>st</sup> operation is Push, it pushes z onto the stack s and continues evaluating the remaining ops \*/

# Evaluation of a list of VM operations

Fixpoint `eval_Op` (`s` : list Z) (`ops` : list Op) : option Z :=

match (`ops`, `s`) with

| (`[]`, [`n`]) => Some `n` /\* If the list is empty and there is only one element on the stack, it returns Some n, indicating successful evaluation with the result n \*/

| (`Push` `z` :: `rest`, \_) => eval\_Op (`z` :: `s`) `rest` /\* If the 1<sup>st</sup> operation is Push, it pushes `z` onto the stack `s` and continues evaluating the remaining ops \*/

| (`Add` :: `rest`, `n2` :: `n1` :: `ns`) => eval\_Op (`n1` + `n2` :: `ns`) %Z `rest`

/\* pops the top `n1` and `n2` from the stack, adds them together, and pushes the result back onto the stack \*/

# Evaluation of a list of VM operations

Fixpoint `eval_Op` (s : list Z) (ops : list Op) : option Z :=

match (ops, s) with

| ([], [n]) => Some n /\* If the list is empty and there is only one element on the stack, it returns Some n, indicating successful evaluation with the result n \*/

| (Push z :: rest, \_) => eval\_Op (z :: s) rest /\* If the 1<sup>st</sup> operation is Push, it pushes z onto the stack s and continues evaluating the remaining ops \*/

| (Add :: rest, n2 :: n1 :: ns) => eval\_Op (n1 + n2 :: ns) %Z rest

/\* pops the top n1 and n2 from the stack, adds them together, and pushes the result back onto the stack \*/

| (OpCount :: rest, \_) => eval\_Op (Z.of\_nat (length rest) :: s) rest /\* counts the remaining ops, converts it to a Z and pushes it onto the stack \*/

# Evaluation of a list of VM operations

```
Fixpoint eval_Op (s : list Z) (ops : list Op) : option Z :=
match (ops, s) with
| ([], [n]) => Some n /* If the list is empty and there is only one element on the
stack, it returns Some n, indicating successful evaluation with the result n */
| (Push z :: rest, _) => eval_Op (z :: s) rest /* If the 1st operation is Push, it
pushes z onto the stack s and continues evaluating the remaining ops */
| (Add :: rest, n2 :: n1 :: ns) => eval_Op (n1 + n2 :: ns)%Z rest
/* pops the top n1 and n2 from the stack, adds them together, and pushes
the result back onto the stack */
| (OpCount :: rest, _) => eval_Op (Z.of_nat (length rest) :: s) rest /* counts the
remaining ops, converts it to a Z and pushes it onto the stack */
| _ => None end. /* if none matches, returns none */
```

# Evaluation of a list of VM operations

## Example

Starting from an empty stack and with the list of operations [Push 1; Push 2; Add], eval\_Op will return 3

The stack evolves as follows

[ ]

[1]

[2 1]

[3]

The evaluation succeeds, and returns 'Some 3'



# Compilation

We compile IMP arithmetic expressions represented by the Expr type into a list of operations (Op) for our stack-based VM

Fixpoint `compile_Expr` (e : Expr) : list Op :=

match e with

| `Num` n => [Push n] /\* If e is a number n, it creates a singleton list containing Push n. This operation pushes n onto the stack.\*/

# Compilation

We compile IMP arithmetic expressions represented by the Expr type into a list of operations (Op) for our stack-based VM

Fixpoint `compile_Expr` (e : Expr) : list Op :=

match e with

| `Num n` => [Push n] /\* If e is a number n, it creates a singleton list containing Push n. This operation pushes n onto the stack.\*/

| `Plus e1 e2` => compile\_Expr e1 ++ compile\_Expr e2 ++ [Add]

/\* If e is an addition, it recursively compiles e1 and e2 into lists of operations, concatenates these lists, and appends an Add operation to the end.\*/

end.

# Compilation

## Example

`compile_Expr (Plus (Num 1) (Num 2)) =>`

# Compilation

## Example

`compile_Expr (Plus (Num 1) (Num 2)) =>`

`compile_Expr (Num 1) ++ compile_Expr (Num 2) ++ [Add] =>`

# Compilation

## Example

```
compile_Expr (Plus (Num 1) (Num 2)) =>  
compile_Expr (Num 1) ++ compile_Expr (Num 2) ++ [Add] =>  
[Push 1] ++ [Push 2] ++ [Add] =
```

# Compilation

## Example

```
compile_Expr (Plus (Num 1) (Num 2)) =>  
compile_Expr (Num 1) ++ compile_Expr (Num 2) ++ [Add] =>  
[Push 1] ++ [Push 2] ++ [Add] =  
[Push 1; Push 2; Add]
```

# Correctness: lemma

**Lemma**  $\text{eval\_step} : \text{forall } a : \text{Expr}, \text{forall } s : \text{list } Z, \text{forall } xs : \text{list } \text{Op},$   
 $\text{eval\_Op } s (\text{compile\_Expr } a ++ xs) = \text{eval\_Op } (\text{eval\_Expr } a :: s) xs.$

# Correctness: lemma

**Lemma**  $\text{eval\_step} : \text{forall } a : \text{Expr}, \text{forall } s : \text{list } Z, \text{forall } xs : \text{list } \text{Op},$   
 $\text{eval\_Op } s (\text{compile\_Expr } a ++ xs) = \text{eval\_Op } (\text{eval\_Expr } a :: s) xs.$

Evaluating a list of operations produced by compiling an expression  $a$  and concatenating it with another list of operations  $xs$

is equivalent to

evaluating the original expression  $a$  and pushing its result onto the stack  $s$ , and then evaluating the list of operations  $xs$

It confirms that the compilation process preserves the semantics of the original expression



# Compiler correctness

Theorem `compiler_correctness` :

forall  $a : \text{Expr}$ ,

$\text{eval\_Op } [] (\text{compile\_Expr } a) = \text{Some } (\text{eval\_Expr } a).$

If you

- compile an expression  $a$  into a list of operations and then
  - evaluate these operations starting with an empty stack,
- you obtain the same result as directly evaluating the original expression  $a$

# Equivalences

Equivalences follow from evaluation

- Definition `equiv_Expr` ( $e1\ e2 : \text{Expr}$ ) :  $\text{Prop} := \text{eval\_Expr } e1 = \text{eval\_Expr } e2$ .

$e1$  and  $e2$  are equivalent if their evaluations ( $\text{eval\_Expr } e1$  and  $\text{eval\_Expr } e2$ , respectively) produce the same result, i.e.,  $e1 \simeq e2$

- Definition `equiv_Op` ( $p1\ p2 : \text{list Op}$ ) :  $\text{Prop} := \text{eval\_Op } []\ p1 = \text{eval\_Op } []\ p2$ .

$p1$  and  $p2$  are equivalent if their evaluations, when starting with an empty stack ( $[]$ ), produce the same result, i.e.,  $p1 \simeq p2$

# Source contexts

Inductive `ExprCtxt` : Set :=

- | `Hole` : `ExprCtxt` /\*hole in the expression where a subexpression can be inserted \*/
- | `Plus1` : `ExprCtxt` -> `Expr` -> `ExprCtxt` /\* context where the right operand of an addition operation is missing\*/
- | `Plus2` : `Expr` -> `ExprCtxt` -> `ExprCtxt` /\*context where the left operand of an addition operation is missing \*/

Expression contexts are used to represent partially completed expressions where certain subexpressions are left as "holes" to be filled in later

# Target contexts

For our stack machine, partial programs are much easier

A program is just a list of Op, therefore any program can be extended by adding new Ops on either end (or inserting in the middle)

# Filling the holes

Fixpoint `link_Expr` (c : ExprCtxt) (e : Expr) : Expr :=

match c with

| `Hole` => e /\*If the context is a Hole, it returns the expression e. This effectively fills the hole with the expression e\*/

| `Plus1` c' e' => Plus (link\_Expr c' e) e' /\* e needs to be inserted as the right operand of an addition operation, with c' representing the left operand. \*/

| `Plus2` e' c' => Plus e' (link\_Expr c' e) /\* e needs to be inserted as the left operand of an addition operation, with c' representing the right operand. \*/

end.

The function `link_Expr` constructs complete expressions by linking an expression e into an expression context c

# Context equivalences

At source level

Definition `ctxtequiv_Expr` ( $e1\ e2 : \text{Expr}$ ) : Prop :=  
forall  $c : \text{ExprCtxt}$ ,  $\text{eval\_Expr} (\text{link\_Expr } c\ e1) = \text{eval\_Expr} (\text{link\_Expr } c\ e2)$ .

$$e1 \simeq_{ctx} e2$$

At target level

Definition `ctxtequiv_Op` ( $p1\ p2 : \text{list Op}$ ) : Prop := forall  $c1\ c2 : \text{list Op}$ ,  
 $\text{eval\_Op } []\ (c1 ++ p1 ++ c2) = \text{eval\_Op } []\ (c1 ++ p2 ++ c2)$ .

$$p1 \simeq_{ctx} p2$$

# Reflection and preservations

We have now all the ingredients to prove

- Reflection and
- Preservation

# Reflection

$$\forall e1, e2. e1 \simeq_{ctx} e2 \Leftarrow \llbracket e1 \rrbracket \simeq_{ctx} \llbracket e2 \rrbracket$$

Lemma equivalence\_reflection :

forall e1 e2 : Expr, forall p1 p2 : list Op,  
forall comp1 : compile\_Expr e1 = p1,  
forall comp2 : compile\_Expr e2 = p2,  
forall eqtarget : ctxtequiv\_Op p1 p2, ctxtequiv\_Expr e1 e2.

If two compiled expressions are **contextually equivalent** at the target level,  
then the original expressions will also be **contextually equivalent** at the  
source level.



# Preservation

$$\forall e1, e2. e1 \simeq_{ctx} e2 \Rightarrow \llbracket e1 \rrbracket \simeq_{ctx} \llbracket e2 \rrbracket$$

Lemma equivalence\_preservation :

forall e1 e2 : Expr, forall p1 p2 : list Op,

forall comp1 : compile\_Expr e1 = p1,

forall comp2 : compile\_Expr e2 = p2,

forall eqsource : ctxtequiv\_Expr e1 e2, ctxtequiv\_Op p1 p2.

If two expressions are **contextually equivalent** at the source level, then their compiled forms will also be **contextually equivalent** at the target level.

Easy to write, but...

# Preservation

$$\forall e1, e2. e1 \simeq_{ctx} e2 \Rightarrow \llbracket e1 \rrbracket \simeq_{ctx} \llbracket e2 \rrbracket$$

Lemma equivalence\_preservation :

forall e1 e2 : Expr, forall p1 p2 : list Op,

forall comp1 : compile\_Expr e1 = p1,

forall comp2 : compile\_Expr e2 = p2,

forall eqsource : ctxtequiv\_Expr e1 e2, ctxtequiv\_Op p1 p2.

If two expressions are **contextually equivalent** at the source level, then their compiled forms will also be **contextually equivalent** at the target level.

Easy to write, but **it is not provable, because it is not true**

# Counterexample

At source level

$$\text{Plus (Num 1) (Num 2)} \simeq_{ctx} \text{(Num 3)}$$

because they have the same evaluation, no matter which is the context

`src_equiv : ctxtequiv_Expr (Plus (Num 1) (Num 2)) (Num 3).`

# Counterexample

At target level, the compiled counterparts are not contextually equivalent, because there exists at least a context that leads to different evaluation results

```
eval_Op [] (OpCount :: compile_Expr (Plus (Num 1) (Num 1)) ++ [Add]) <>  
eval_Op [] (OpCount :: compile_Expr (Num 2) ++ [Add]).
```

**Why?**

# Counterexample (cont.)

```
eval_Op [] (OpCount :: compile_Expr (Plus (Num 1) (Num 1)) ++ [Add]) <>  
eval_Op [] (OpCount :: compile_Expr (Num 2) ++ [Add]).
```

If we put **OpCount** followed by Add instruction afterwards, the result will be:

- the value and
- the number of instructions it took to compute it, that is different
  - for `OpCount :: compile_Expr (Plus (Num 1) (Num 1)) ++ [Add]` is [6].
  - for `OpCount :: compile_Expr (Num 2) ++ [Add]` is [4].
- Note that OpCount context at target level does not have a corresponding context (obtainable by back-translation) at source level

# Counterexample

This is an example of common situation

There could be cases where  $e_1$  and  $e_2$  are contextually equivalent as expressions but their compiled forms  $p_1$  and  $p_2$  might not be equivalent as operations due to

- the structure of the compilation process or
- how the operations interact with the stack machine

# Any solution?

To ensure the correctness of this theorem, we would need:

- **Strong Equivalence Guarantees:**

Prove that `compile_Expr` preserves not just the final results but the entire behavior of expressions in all contexts

- **Correct Handling of Stack State:**

Ensure that operations resulting from `compile_Expr` maintain equivalent stack states across all possible contexts

# Any solution?

We need to somehow protect the compiled code from having the equivalences disrupted

Here, for instance, we might put some flag on instructions that meant that they should not be counted, and OpCount would just not return anything if it saw any of those (or would count them as 0).

Alternately, we might give our target language a **type system** that is able to rule out linking with code that uses the OpCount instruction, or perhaps restricts how it can be used



# To keep it simple

For the sake of simplicity, we opt here for an extreme solution:

- rather than a list, we allow one Op before and one Op after our compiled program, neither of which can be OpCount
- the resulting program to be well-formed (i.e., no errors, only one number on stack at end), so
  - either there should be nothing before and after,
  - or there is a Push n before and either Add or Sub after
  - no other combination of Op before or after will fulfill our requirement

# To keep it simple

We highly restrict what we can link with: rather than a list, we allow one Op before and one Op after our compiled program, neither of which OpCount

Inductive **OpCtxt** : Set :=  
| **PushAdd** :  $\mathbb{Z} \rightarrow \text{OpCtxt}$   
| **Empty** : OpCtxt.

Definition **link\_Op** (c : OpCtxt) (p : list Op) : list Op :=  
match c with  
| **PushAdd** n => Push n :: p ++ [Add]  
| **Empty** => p end.

# To keep it simple

We need to redefine contextual equivalence for our target language, only permitting these new contexts

Definition `ctxtequiv_Op` (p1 p2 : list Op) : Prop :=  
forall c : OpCtxt, eval\_Op [] (link\_Op c p1) = eval\_Op [] (link\_Op c p2).

Now our compiler, when linked against these restricted contexts, can be proved fully abstract

# Back-translation

To complete the proof, we rely upon a **back-translation** from target contexts to source context, where target contexts are restricted as described above:  
from OpCtxt to ExprCtxt

Definition **backtranslate** (**c** : OpCtxt) : ExprCtxt :=

match c with

| **PushAdd** n => Plus2 (Num n) Hole

*/\* the context involving pushing n and the add is back-translated into a context where the expression (Num n) + Hole is formed \*/*

# Back-translation

To complete the proof, we rely upon a **back-translation** from target contexts to source context, where target contexts are restricted as described above:  
from OpCtxt to ExprCtxt

Definition **backtranslate** (**c** : OpCtxt) : ExprCtxt :=

match c with

| **PushAdd** n => Plus2 (Num n) Hole

*/\* the context involving pushing n and the add is back-translated into a context where the expression (Num n) + Hole is formed \*/*

| **Empty** => Hole end. */\*an empty context is back-translated in a context a context with no additional expressions or operations \*/*

# Back-translation lemma

Lemma `back_translation_equiv` : forall c : OpCtxt, forall p : list Op, forall e : Expr, forall c' : ExprCtxt, compile\_Expr e = p -> backtranslate c = c' -> eval\_Op [] (link\_Op c p) = Some (eval\_Expr (link\_Expr c' e)).

The lemma aims to prove that evaluating the **operational context** applied to the compiled expression yields the same result as evaluating the **expression context** applied to the original expression

# Equivalence preservation

Now, we can prove equivalence preservation directly.

Lemma `equivalence_preservation` :

forall e1 e2 : Expr, forall p1 p2 : list Op,  
forall comp1 : compile\_Expr e1 = p1,  
forall comp2 : compile\_Expr e2 = p2,  
forall eqsource : ctxtequiv\_Expr e1 e2, ctxtequiv\_Op p1 p2.

The theorem is used to prove that

$\forall e1, e2. e1 \simeq_{ctx} e2 \Rightarrow \llbracket e1 \rrbracket \simeq_{ctx} \llbracket e2 \rrbracket$

# Proving equivalence

This is obviously a very tiny language and a very restrictive linker that only allowed very restrictive contexts, but the general shape of the proof is the same as that used in more realistic languages



# Secure compilation

Must ensure that any P we link with behaves like some source context

- Add target features to the source language **Bad!**
- Dynamics checks: catch badly behaved code in the act **Expensive**
- Static checks: rule out badly behaved code in the first place **Verification**

# Secure compilation: recap

- Preserving security by preserving equivalence
- Different compilation targets and threat models
  - is the target language typed or untyped?
  - what observations can the attacker make?
- Different ways of enforcing secure compilation
  - static checking
  - dynamic checking (e.g., runtime monitoring, hardware enforcement)
- Proof techniques: "back-translating" target attackers to source

# Bibliography

- Marco Patrignani, Amal Ahmed, Dave Clarke. *Formal Approaches to Secure Compilation. A Survey of Fully Abstract Compilation and Related Work*. ACM Computing surveys, 2019.
- <https://dbp.io/essays/2018-04-19-how-to-prove-a-compiler-fully-abstract.html>

**End**